

Pipelining Wrap-Up

Lecture 14

November 16th, 2017

Jae W. Lee (jaewlee@snu.ac.kr)

Computer Science and Engineering

Seoul National University

Slide credits: [CS:APP3e] slides from CMU; [COD5e] slides from Elsevier Inc.

Today

Textbook: [CS:APP3e] 4.5.6, 4.5.9, and 5.7

- **Wrap-Up of PIPE Design**

- Exceptional conditions
- Performance analysis

- **Modern High-Performance Processors**

- Out-of-order execution


Exceptions

- Conditions under which processor cannot continue normal operation

■ Causes

- Halt instruction (Current)
- Bad address for instruction or data (Previous)
- Invalid instruction (Previous)

■ Typical Desired Action

- Complete some instructions
 - Either current or previous (depends on exception type) 
- Discard others
- Call exception handler
 - Like an unexpected procedure call

■ Our Implementation

- Halt when instruction causes exception

Exception Examples

■ Detect in Fetch Stage

`jmp $-1` `# Invalid jump target`

`.byte 0xFF` `# Invalid instruction code`

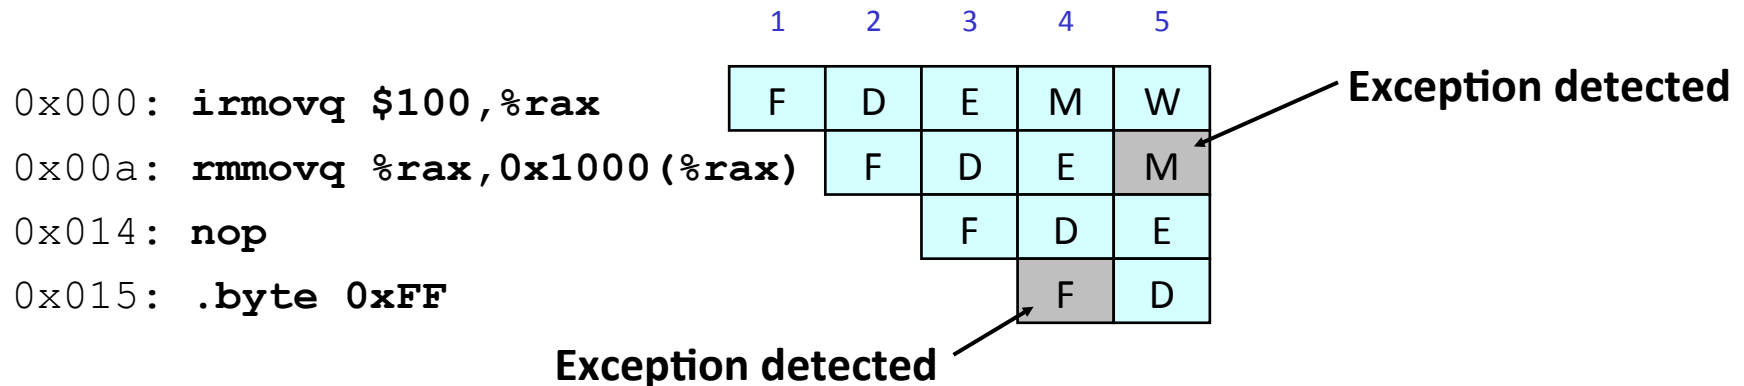
`halt` `# Halt instruction`

■ Detect in Memory Stage

`irmovq $100,%rax`
`rmmovq %rax,0x10000(%rax) # invalid address`

Exceptions in Pipeline Processor #1

```
# demo-excl1.y
irmovq $100,%rax
rmmovq %rax,0x10000(%rax) # Invalid address
nop
.byte 0xFF                # Invalid instruction code
```



■ Desired Behavior

- `rmmovq` should cause exception
- Following instructions should have no effect on processor state

Exceptions in Pipeline Processor #2

```
# demo-exc2.ys
```

```
0x000:    xorq %rax,%rax    # Set condition codes
```

```
0x002:    jne t            # Not taken
```

```
0x00b:    irmovq $1,%rax
```

```
0x015:    irmovq $2,%rdx
```

```
0x01f:    halt
```

```
0x020: t: .byte 0xFF      # Target
```


```
0x000:    xorq %rax,%rax
```

```
0x002:    jne t
```

```
0x020: t: .byte 0xFF
```

```
0x???: (I'm lost!)
```

```
0x00b:    irmovq $1,%rax
```

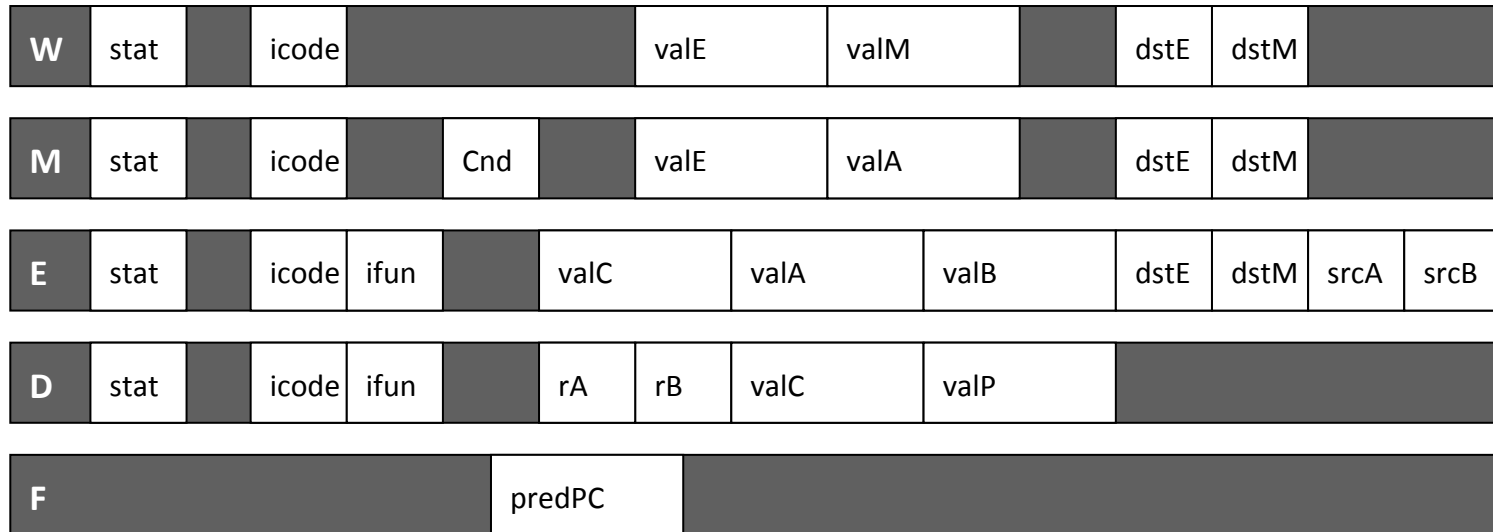


Exception detected

■ Desired Behavior

- No exception should occur

Maintaining Exception Ordering



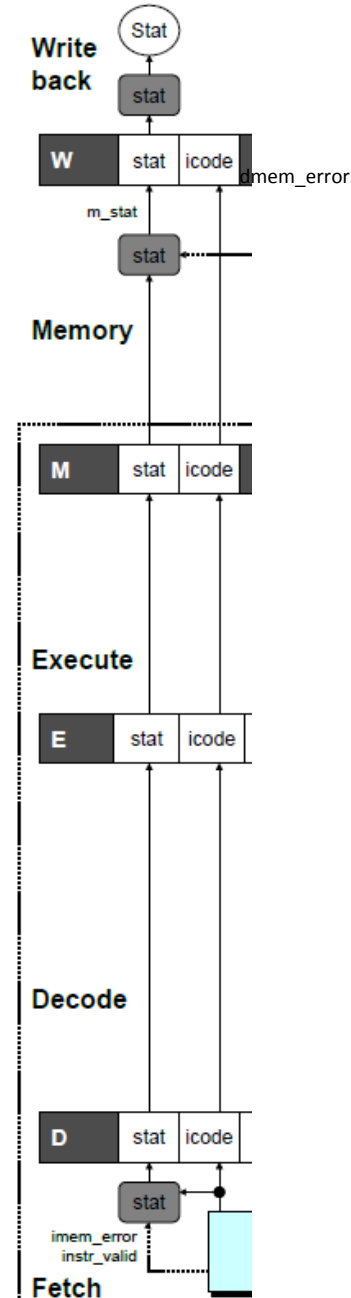
- Add status field to pipeline registers
- Fetch stage sets to either “AOK,” “ADR” (when bad fetch address), “HLT” (halt instruction) or “INS” (illegal instruction)
- Decode & execute pass values through
- Memory either passes through or sets to “ADR”
- Exception triggered only when instruction hits write back

■ Fetch Stage

```
int f_stat = [
    imem_error: SADR;
    !instr_valid : SINS;
    f_icode == IHALT : SHLT;
    1 : SAOK;
];
```

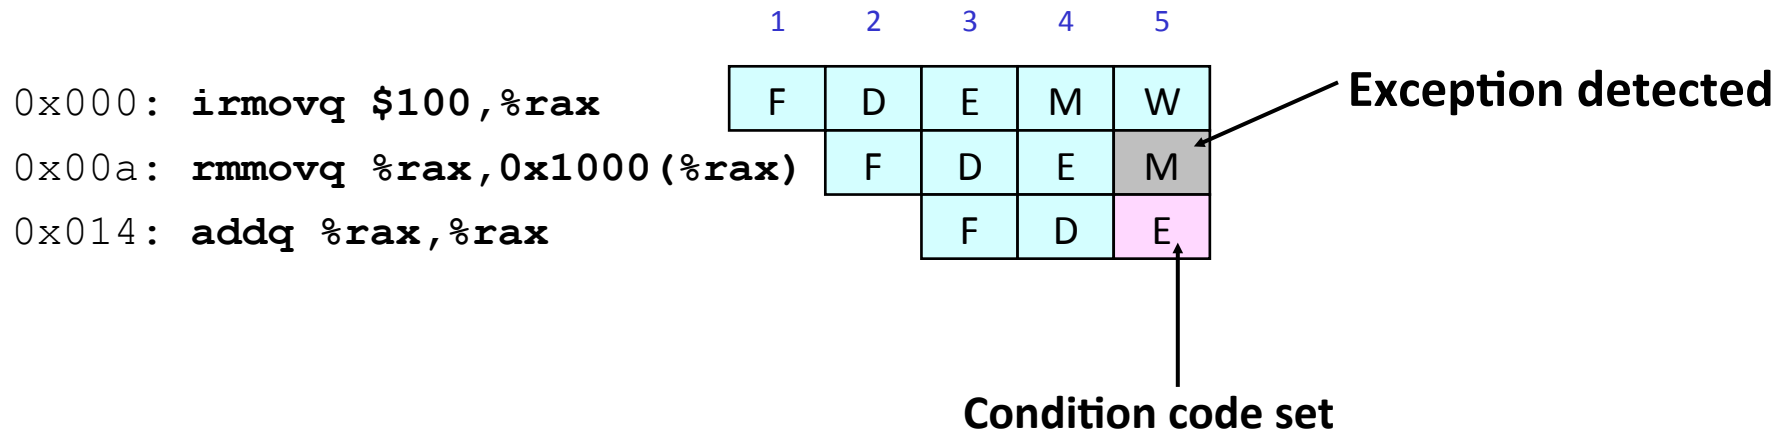
```
# Update the status
int m_stat = [
    dmem_error : SADR;
    1 : M_stat;
];
```

```
int Stat = [
    # SBUB in earlier stages indicates bubble
    W_stat == SBUB : SAOK;
    1 : W_stat;
];
```



Side Effects in Pipeline Processor

```
# demo-exc3.js
irmovq $100,%rax
rmmovq %rax,0x10000(%rax) # invalid address
addq %rax,%rax             # Sets condition codes
```



■ Desired Behavior

- `rmmovq` should cause exception
- No following instruction should have any effect

Avoiding Side Effects

- **Presence of Exception Should Disable State Update**
 - Invalid instructions are converted to pipeline bubbles
 - Except have stat indicating exception status
 - Data memory will not write to invalid address
 - Prevent invalid update of condition codes
 - Detect exception in memory stage
 - Disable condition code setting in execute
 - Must happen in same clock cycle
 - Handling exception in final stages
 - When detect exception in memory stage
 - Start injecting bubbles into memory stage on next cycle
 - When detect exception in write-back stage
 - Stall excepting instruction

Control Logic for State Changes

■ Setting Condition Codes

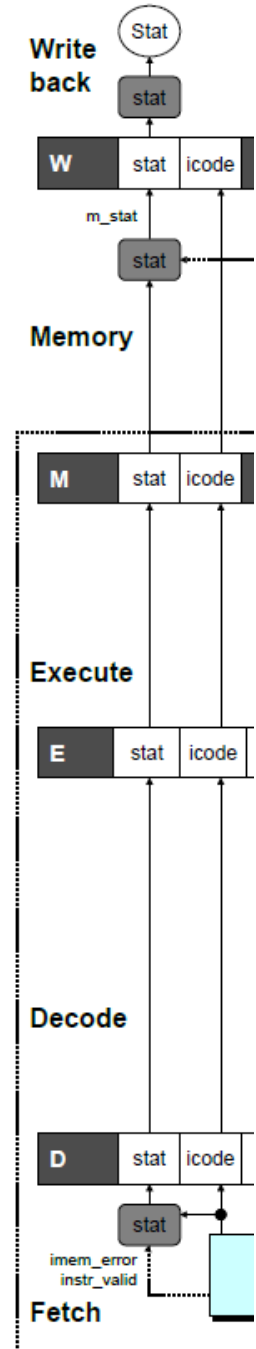
```
# Should the condition codes be updated?
bool set_cc = E_icode == IOPQ &&
    # State changes only during normal operation
    !m_stat in { SADR, SINS, SHLT }
    && !W_stat in { SADR, SINS, SHLT };
```

■ Stage Control

- Also controls updating of memory

```
# Start injecting bubbles as soon as exception passes
through memory stage
bool M_bubble = m_stat in { SADR, SINS, SHLT }
    || W_stat in { SADR, SINS, SHLT };
```

```
# Stall pipeline register W when exception encountered
bool W_stall = W_stat in { SADR, SINS, SHLT };
```



Rest of Real-Life Exception Handling

■ Call Exception Handler

- Push PC onto stack
 - Either PC of faulting instruction or of next instruction
 - Usually pass through pipeline along with exception status
- Jump to handler address
 - Usually fixed address
 - Defined as part of ISA

Performance Metrics

■ Clock rate

- Measured in Gigahertz
- Function of stage partitioning and circuit design
 - Keep amount of work per stage small

■ Rate at which instructions executed

- CPI: cycles per instruction
- On average, how many clock cycles does each instruction require?
- Function of pipeline design and benchmark programs
 - E.g., how frequently are branches mispredicted?

CPI for PIPE

■ CPI \approx 1.0

- Fetch instruction each clock cycle
- Effectively process new instruction almost every cycle
 - Although each individual instruction has latency of 5 cycles

■ CPI > 1.0

- Sometimes must stall or cancel branches

■ Computing CPI

- C clock cycles
- I instructions executed to completion
- B bubbles injected ($C = I + B$)

$$\text{CPI} = C/I = (I+B)/I = 1.0 + B/I$$

- Factor B/I represents average penalty due to bubbles

CPI for PIPE (Cont.)

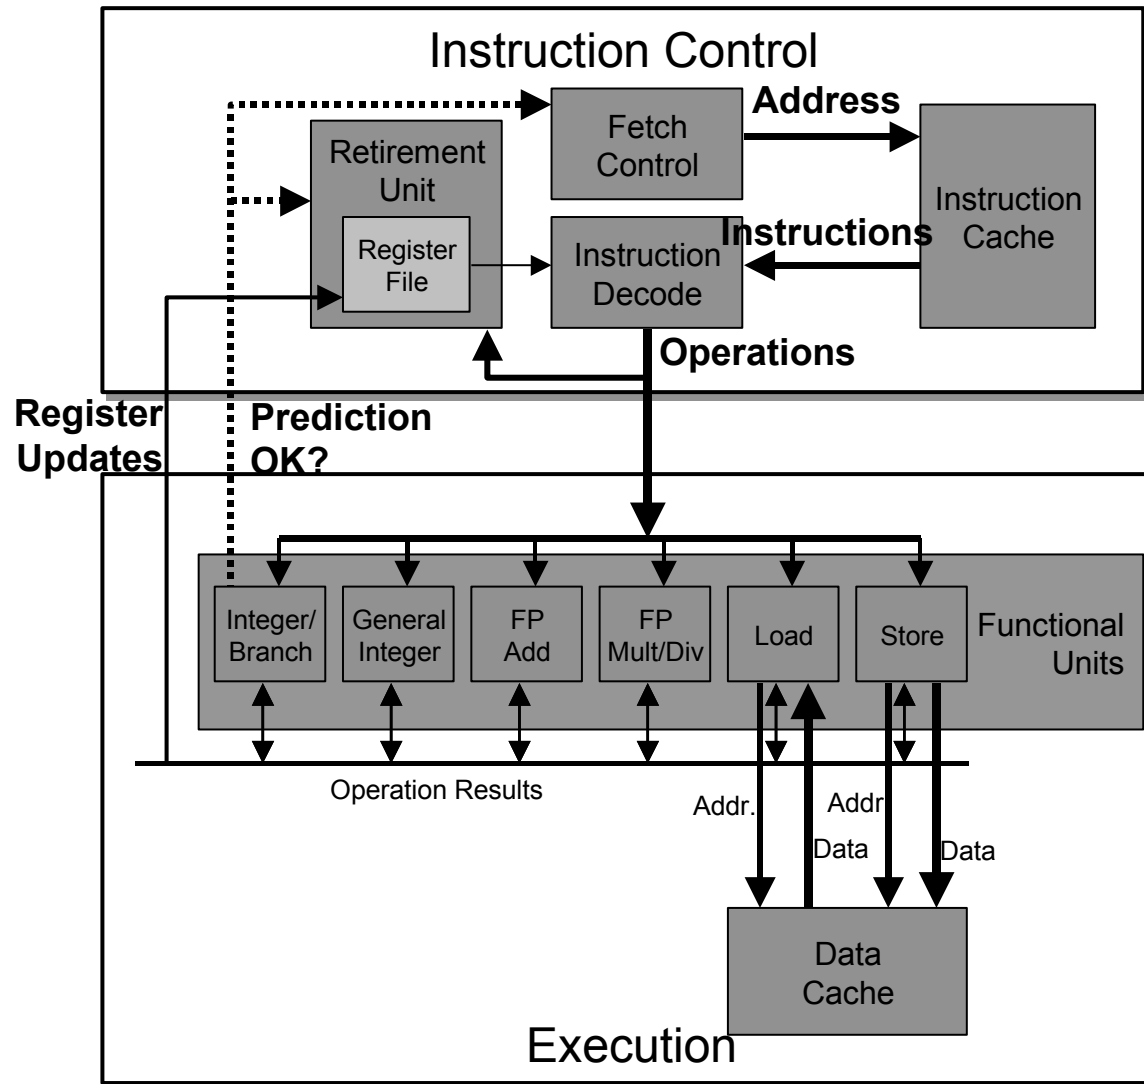
$$B/I = LP + MP + RP$$

Typical Values

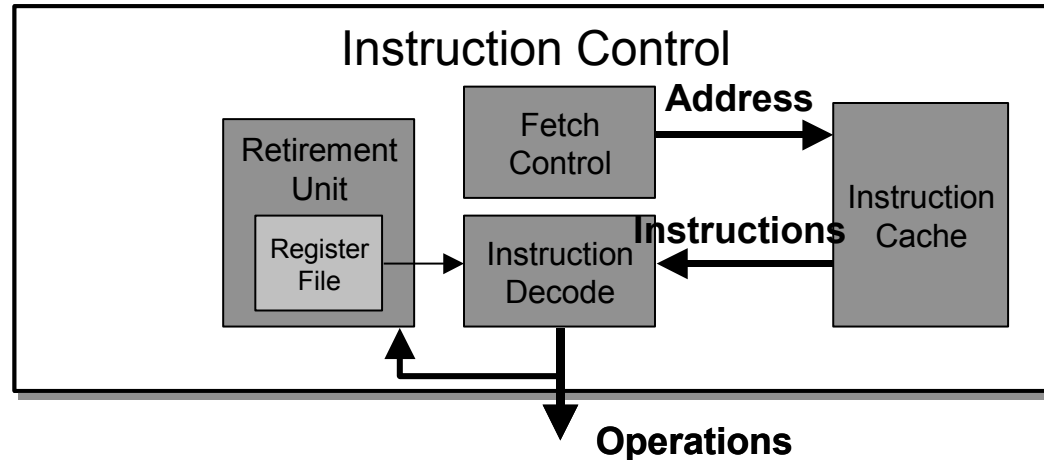
- LP: Penalty due to load/use hazard stalling
 - Fraction of instructions that are loads 0.25
 - Fraction of load instructions requiring stall 0.20
 - Number of bubbles injected each time 1
$$\Rightarrow LP = 0.25 * 0.20 * 1 = 0.05$$
- MP: Penalty due to mispredicted branches
 - Fraction of instructions that are cond. jumps 0.20
 - Fraction of cond. jumps mispredicted 0.40
 - Number of bubbles injected each time 2
$$\Rightarrow MP = 0.20 * 0.40 * 2 = 0.16$$
- RP: Penalty due to `ret` instructions
 - Fraction of instructions that are returns 0.02
 - Number of bubbles injected each time 3
$$\Rightarrow RP = 0.02 * 3 = 0.06$$
- Net effect of penalties $0.05 + 0.16 + 0.06 = 0.27$

$$\Rightarrow CPI = 1.27 \text{ (Not bad!)}$$

Modern CPU Design



Instruction Control



■ Grabs Instruction Bytes From Memory

- Based on Current PC + Predicted Targets for Predicted Branches
- Hardware dynamically guesses whether branches taken/not taken and (possibly) branch target

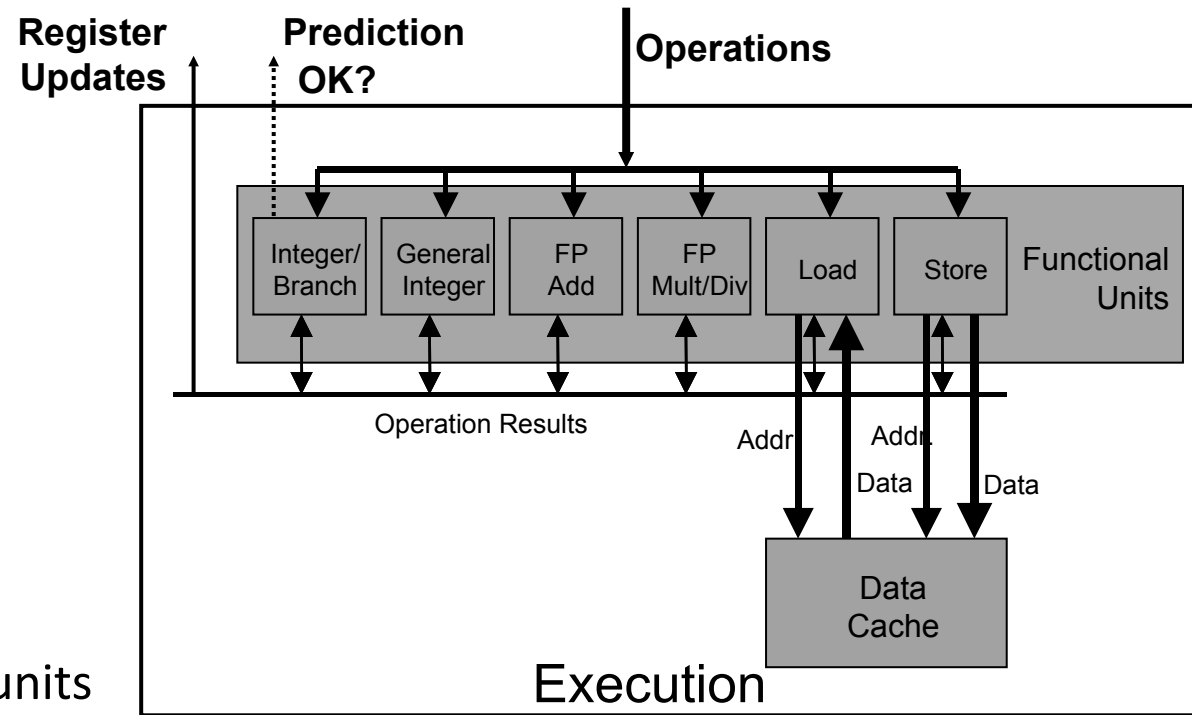
■ Translates Instructions Into *Operations*

- Primitive steps required to perform instruction
- Typical instruction requires 1–3 operations

■ Converts Register References Into *Tags*

- Abstract identifier linking destination of one operation with sources of later operations

Execution Units



- Multiple functional units
 - Each can operate independently
- Operations performed as soon as operands available
 - Not necessarily in program order
 - Within limits of functional units
- Control logic
 - Ensures behavior equivalent to sequential program execution

CPU Capabilities of Intel Haswell

■ Multiple Instructions Can Execute in Parallel

- 2 load
- 1 store
- 4 integer
- 2 FP multiply
- 1 FP add / divide

■ Some Instructions Take > 1 Cycle, but Can be Pipelined

■ Instruction	Latency	Cycles/Issue
■ Load / Store	4	1
■ Integer Multiply	3	1
■ Integer Divide	3—30	3—30
■ Double/Single FP Multiply	5	1
■ Double/Single FP Add	3	1
■ Double/Single FP Divide	10—15	6—11

Haswell Operation

- **Translates instructions dynamically into “Uops”**
 - ~118 bits wide
 - Holds operation, two sources, and destination
- **Executes Uops with “Out of Order” engine**
 - Uop executed when
 - Operands available
 - Functional unit available
 - Execution controlled by “Reservation Stations”
 - Keeps track of data dependencies between uops
 - Allocates resources

High-Performance Branch Prediction

■ Critical to Performance

- Typically 11–15 cycle penalty for misprediction

■ Branch Target Buffer

- 512 entries
- 4 bits of history
- Adaptive algorithm
 - Can recognize repeated patterns, e.g., alternating taken–not taken


■ Handling BTB misses

- Detect in ~cycle 6
- Predict taken for negative offset, not taken for positive
 - Loops vs. conditionals

Example Branch Prediction

■ Branch History

- Encode information about prior history of branch instructions
- Predict whether or not branch will be taken



■ State Machine

- Each time branch taken, transition to right
- When not taken, transition to left
- Predict branch taken when in state **Yes!** or **Yes?**

Processor Summary

■ Design Technique

- Create uniform framework for all instructions
 - Want to share hardware among instructions
- Connect standard logic blocks with bits of control logic

■ Operation

- State held in memories and clocked registers
- Computation done by combinational logic
- Clocking of registers/memories sufficient to control overall behavior

■ Enhancing Performance

- Pipelining increases throughput and improves resource utilization
- Must make sure to maintain ISA behavior