

Synchronization (1)



April 2, 2018
Byung-Gon Chun

Acknowledgments. Some slides and/or picture in the following are adapted from UW, Columbia, and UC Berkeley class slides

Why Synchronization?

- If a program has independent threads that operate on completely separate subsets of memory, we can reason about each thread separately.
- But, most multi-threaded programs have both per-thread state and shared state. Cooperating threads read and write shared state

=> Writing correct multi-threaded programs becomes much more difficult.
- The sequential model of reasoning does not work in programs with cooperating threads

Why Does Not The Sequential Model of Reasoning Work?

- Program execution depends on the possible interleavings of threads' access to shared state
 - Two threads write to the same variable; which one should win?
- Program execution can be nondeterministic
 - Different runs of the same program may produce different results
 - Heisenbugs: bugs that disappear or change behavior when you try to examine them
- Compilers and processor hardware can reorder instructions

Question: Can this panic?

Thread 1

```
p = someComputation();  
pInitialized = true;
```

Thread 2

```
while (!pInitialized)  
    ;  
q = anotherComputation (p);  
if (q != anotherComputation(p))  
    panic
```

Roadmap

- Synchronization Challenges
- Structuring Shared Objects
- Locks: Mutual Exclusion
- Condition Variables: Waiting for a Change
- Designing and Implementing Shared Objects
- Case Studies
- Implementing Synchronization Primitives
- Semaphore
- Linux kernel synchronization
- Multiprocessor Lock Performance
- Lock Design Patterns
- Lock Contention
- Multi-Object Atomicity
- Deadlock
- Non-Blocking Synchronization

Challenges

Race condition: output of a concurrent program depends on the order of operations between threads

■ Can be very bad

- “non-deterministic:” don’t know what the output will be, and it is likely to be different across runs
- Hard to detect: too many possible schedules
- Hard to debug: “heisenbug,” debugging changes timing so hides bugs (vs “bohr bug”)

Simple Cooperating-Threads Program 1

- A program with two threads that do the following

Thread A

$x = 1;$

Thread B

$x = 2;$

What are the possible final values of x ?

Simple Cooperating-Threads Program 2

- A program with two threads that do the following. Initially $y = 12$.

Thread A

$x = y + 1;$

Thread B

$y = y * 2;$

What are the possible final values of x ?

Simple Cooperating-Threads Program 3

- A program with two threads that do the following. Initially $x = 0$.

Thread A

$x = x + 1;$

Thread B

$x = x + 2;$

What are the possible final values of x ?

Banking Example

```
int balance = 0;
int main() {
    pthread_t t1, t2;
    pthread_create(&t1, NULL, deposit, (void*)1);
    pthread_create(&t2, NULL, withdraw, (void*)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("all done: balance = %d\n", balance);
    return 0;
}

void* deposit(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i)
        ++ balance;
}

void* withdraw(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i)
        -- balance;
}
```

Results of the banking example

```
$ gcc -Wall -lpthread -o bank bank.c
```

```
$ bank
```

```
all done: balance = 0
```

```
$ bank
```

```
all done: balance = 140020
```

```
$ bank
```

```
all done: balance = -94304
```

```
$ bank
```

```
all done: balance = -191009
```

Why?

A closer look at the banking example

```
$ objdump -d bank
```

```
...
```

```
08048464 <deposit>:
```

```
...
```

```
8048473: a1 80 97 04 08
```

```
8048478: 83 c0 01
```

```
804847b: a3 80 97 04 08
```

```
...
```

```
0804849b <withdraw>:
```

```
...
```

```
80484aa: a1 80 97 04 08
```

```
80484af: 83 e8 01
```

```
80484b2: a3 80 97 04 08
```

```
...
```

```
// ++ balance
```

```
mov 0x8049780,%eax
```

```
add $0x1,%eax
```

```
mov %eax,0x8049780
```

```
// -- balance
```

```
mov 0x8049780,%eax
```

```
sub $0x1,%eax
```

```
mov %eax,0x8049780
```

One possible schedule

CPU 0

mov 0x8049780,%eax

eax: 0

add \$0x1,%eax

eax: 1

mov %eax,0x8049780

balance: 1

CPU 1

mov 0x8049780,%eax

eax: 1

sub \$0x1,%eax

eax: 0

mov %eax,0x8049780

balance: 0

One deposit and one withdraw, balance unchanged. Correct.

Another possible schedule

CPU 0

mov 0x8049780,%eax

eax: 0

add \$0x1,%eax

eax: 1

mov %eax,0x8049780

balance: 1

CPU 1

mov 0x8049780,%eax

eax: 0

sub \$0x1,%eax

eax: -1

mov %eax,0x8049780

balance: -1

One deposit and one withdraw, balance becomes -1. Wrong!

Simple Cooperating-Threads Program 3

- A program with two threads that do the following. Initially $x = 0$.

Thread A

$x = x + 1;$

Thread B

$x = x + 2;$

Interleaving 1

load r1, x
add r2, r1, 1
store x, r2

What are the possible final values of x ?

load r1, x
add r2, r1, 2
store x, r2

Interleaving 2

load r1, x	
	load r1, x
add r2, r1, 1	
	add r2, r1, 2
store x, r2	
	store x, r2

Interleaving 3

load r1, x	
	load r1, x
add r2, r1, 1	
	add r2, r1, 2
	store x, r2
store x, r2	

Sharing a Refrigerator

- Two room mates who share a refrigerator and who make sure the refrigerator is always well stocked with milk.

Too Much Milk Example

	Person A	Person B
12:30	Look in fridge. Out of milk.	
12:35	Leave for store.	
12:40	Arrive at store.	Look in fridge. Out of milk.
12:45	Buy milk.	Leave for store.
12:50	Arrive home, put milk away.	Arrive at store.
12:55		Buy milk.
1:00		Arrive home, put milk away. Oh no!

Simplifying Assumptions to Solve Too Much Milk

- Instructions are executed in exactly the order written.
I.e., neither the compiler nor the architecture reorders instructions

Too Much Milk, Try #1

■ Correctness property

- Someone buys if needed (liveness)
- At most one person buys (safety)

■ Try #1: leave a note

```
if (milk == 0)
    if (note == 0) {
        note = 1; // leave note
        milk++; // buy milk
        note = 0; // remove note
    }
```

Too Much Milk, Try #1

Thread A

```
if (milk == 0) {
```

```
    if (note == 0) {  
        note = 1;  
        milk++;  
        note = 0;  
    }
```

```
}
```

Thread B

```
if (milk==0) {  
    if (note == 0) {  
        note = 1;  
        milk++;  
        note = 0;  
    }  
}
```

Too Much Milk, Try #2

Thread A

```
noteA = 1; // leave note A
if (noteB==0) { // if no note A1
    if (milk==0) // if no milk A2
        milk++ // buy milk A3
    }
noteA = 0; // remove note A
```

Thread B

```
noteB = 1; // leave note B
if (noteA==0) { // if no note B1
    if (milk==0) // if no milk B2
        milk++; // buy milk B3
    } // B4
noteB=0; // remove note B
```

Too Much Milk, Try #3

Thread A

```
noteA=1;    // leave note A
while (noteB == 1) // X: wait for no noteB
    ;        // spin
if (milk==0) // if no milk M
    milk++;  // buy milk
noteA = 0;   // remove note A
```

Thread B

```
noteB = 1;    // leave note B
if (noteA==0) { // Y: if no note A
    if (milk==0) // if no milk
        milk++; // buy milk
}
noteB = 0;    // remove note B
```

Too Much Milk, Try #3

Thread A

```
noteA=1;    // leave note A
while (noteB == 1) // X: wait for no noteB
    ;        // spin
if (milk==0) // if no milk M
    milk++;  // buy milk
noteA = 0;   // remove note A
```

Thread B

```
noteB = 1;    // leave note B
if (noteA==0) { // Y: if no note A
    if (milk==0) // if no milk
        milk++; // buy milk
}
noteB = 0;    // remove note B
```

At Y: if no Note A, safe for B to buy (means A hasn't started yet)
if note A, A is either buying, or waiting for B to quit,
so ok for B to quit

Too Much Milk, Try #3

Thread A

```
noteA=1;    // leave note A
while (noteB == 1) // X: wait for no noteB
    ;        // spin
if (milk==0) // if no milk M
    milk++;  // buy milk
noteA = 0;   // remove note A
```

Thread B

```
noteB = 1;    // leave note B
if (noteA==0) { // Y: if no note A
    if (milk==0) // if no milk
        milk++; // buy milk
}
noteB = 0;    // remove note B
```

At X: if no note B, safe to buy
if note B, don't know. A hangs around. Either:
if B buys, done
if B doesn't buy, A will.

Discussion: Are You Satisfied with the Solution?

- Solution is complicated

- “obvious” code often has bugs

- Solution is inefficient

- While Thread A is waiting, it is busy-waiting and consuming CPU resources

- The solution may fail if the compiler or hardware reorders instructions

- The limitation can be addressed by using memory barriers. This makes reasoning even more difficult

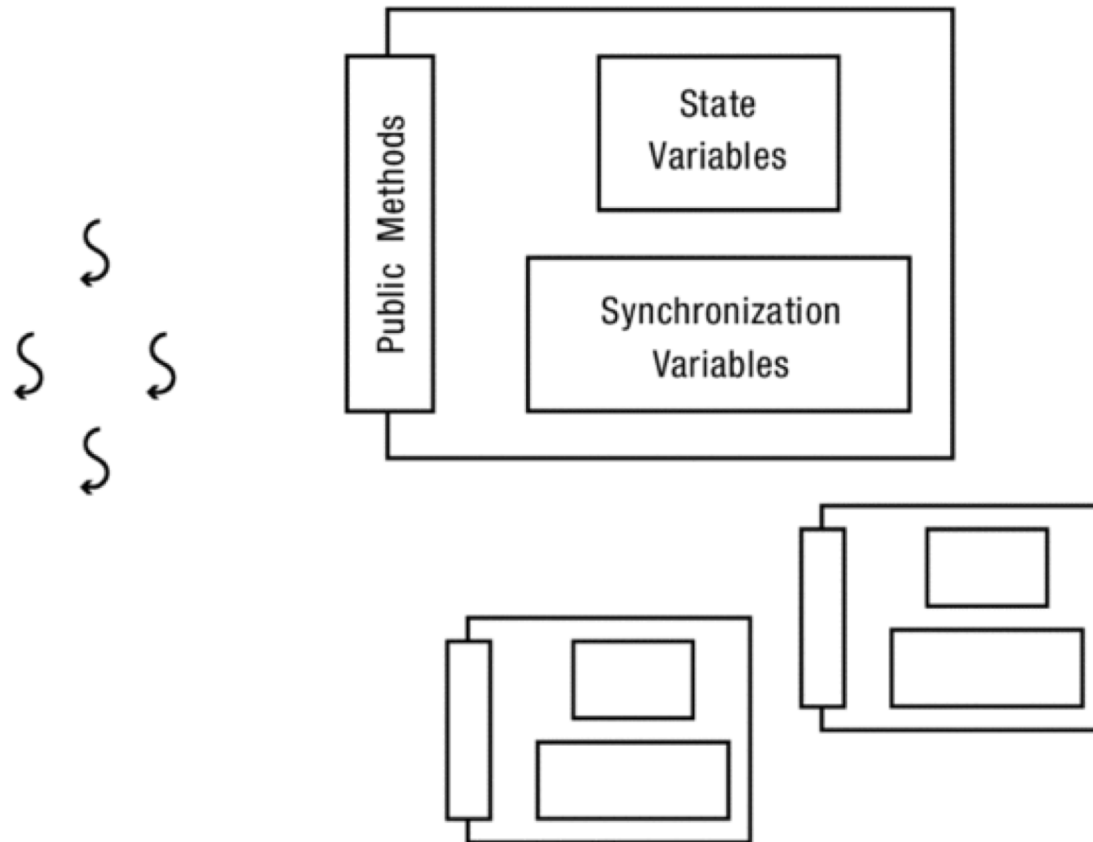
- Generalizing to many threads/processors

- Even more complex: see Peterson’s algorithm

Structuring Shared Objects

Threads

Shared Objects



Implementing Shared Objects

Concurrent Applications

Shared Objects

Bounded Buffer

Barrier

Synchronization Variables

Semaphores

Locks

Condition Variables

Atomic Instructions

Interrupt Disable

Test-and-Set

Hardware

Multiple Processors

Hardware Interrupts

Implementing Shared Objects

Concurrent Applications

Shared Objects

Bounded Buffer

Barrier

Synchronization Variables

Semaphores

Locks

Condition Variables

Atomic Instructions

Interrupt Disable

Test-and-Set

Hardware

Multiple Processors

Hardware Interrupts

Locks

- Lock is a synchronization variable that provides mutual exclusion – when one thread holds a lock, no other thread can hold it
- A program associates each lock with some subset of shared state and requires a thread to hold the lock when accessing that state.
- Mutual exclusion greatly simplifies reasoning about programs because a thread can perform an arbitrary set of operations while holding a lock, and those operations appear to be atomic to other threads

Locks: API

- Two methods: Lock::acquire and Lock::release
- A lock can be in one of two states: BUSY or FREE
- A lock is initially in the FREE state
- Lock::acquire waits until the lock is FREE and then takes it (**atomically** makes the lock BUSY)
- Lock::release makes the lock FREE. If there are pending acquire operations, this state change causes one of them to proceed

Question: Why only Acquire/Release

- Suppose we add a method to a lock, to ask if the lock is free. Suppose it returns true. Is the lock:
 - Free?
 - Busy?
 - Don't know?

Too Much Milk, #4

Locks allow concurrent code to be much simpler:

```
lock.acquire();  
if (milk==0)  
    milk++; // buy milk  
lock.release();
```


Lock Example: Malloc/Free

```
char *malloc (n) {  
    p = allocate memory  
    return p;  
}
```

```
void free(char *p) {  
    put p back on free list  
}
```

How to implement a thread-safe memory allocator?

Lock Example: Malloc/Free

```
char *malloc (n) {  
    heaplock.acquire();  
    p = allocate memory  
    heaplock.release();  
    return p;  
}
```

```
void free(char *p) {  
    heaplock.acquire();  
    put p back on free list  
    heaplock.release();  
}
```

Fixing the Banking Example

```
void* deposit(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i) {
        ++ balance;
    }
}
```

```
void* withdraw(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i) {
        -- balance;
    }
}
```

Fixing the Banking Example

```
pthread_mutex_t lock;
```

```
void* deposit(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i) {
        pthread_mutex_lock(&lock);
        ++ balance;
        pthread_mutex_unlock(&lock);
    }
}
```

```
void* withdraw(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i) {
        pthread_mutex_lock(&lock);
        -- balance;
        pthread_mutex_unlock(&lock);
    }
}
```