Memory-Mapped Files & Virtual Memory

May 23, 2018 Byung-Gon Chun

Acknowlegments. Slides and/or picture in the following are adapted from UW, Columbia, and UC Berkeley slides

Memory Management

- Address Translation
 - Basic concept
 - Flexible
 - Efficient
- Caching
- Virtual memory



Memory-Mapped Files

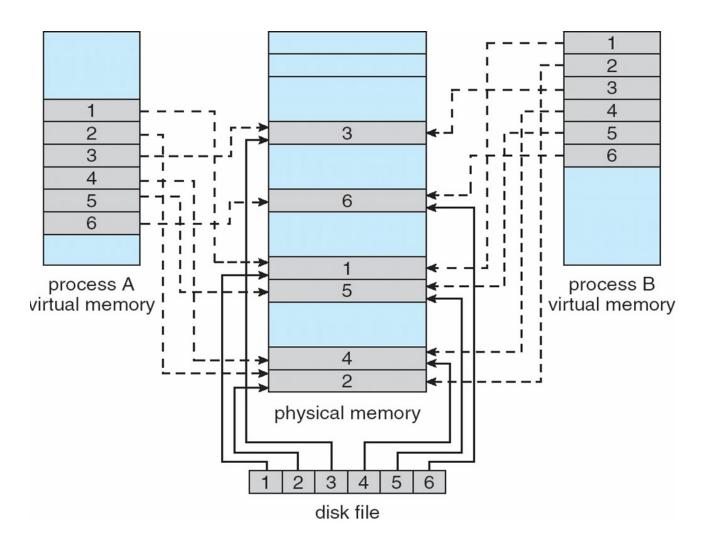
Demand Paging

- With demand paging, applications can access more memory than is physically present on the machine, by using memory pages as a cache for disk blocks.
- When the application accesses a missing memory page, it is transparently brought in from disk.
 - Simpler case of a demand paging: a single, memory-mapped file
 - More complex case: managing multiple processes competing for space in main memory

Models for Application File I/O

- Explicit read/write system calls
 - Data copied to user process using system call
 - Application operates on data
 - Data copied back to kernel using system call
- Memory-mapped files
 - Open file as a memory segment
 - Program uses load/store instructions on segment memory, implicitly operating on the file
 - Page fault if portion of file is not yet in memory
 - Kernel brings missing blocks into memory, restarts process

Memory Mapped File Example



Advantages to Memory-mapped Files

- Programming simplicity, esp for large files
 - Operate directly on file, instead of copy in/copy out
- Zero-copy I/O
 - Data brought from disk directly into page frame
- Pipelining
 - Process can start working before all the pages are populated
- Interprocess communication
 - Shared memory segment vs. temporary file

Implementation

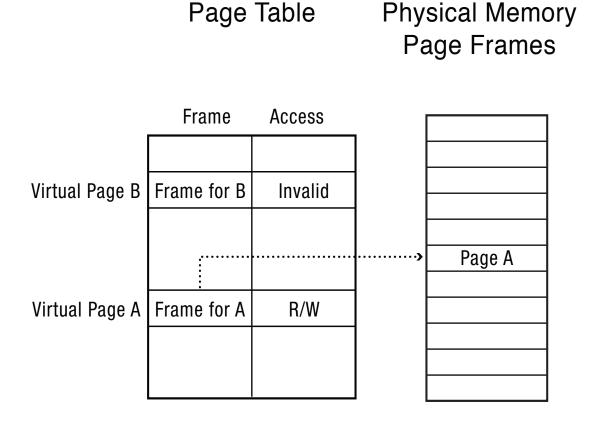
Map the file into a portion of the virtual address space

The kernel initializes a set of page table entries for that region of the virtual address space, setting each entry to invalid.

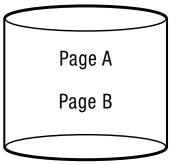
Implementation

- When the process issues an instruction that touches an invalid mapped address, a sequence of events occurs
 - □ TLB miss
 - Page table exception
 - Convert virtual address to file offset
 - □ Disk block read: allocate an empty page frame and issue a disk operation to read the required file block into the allocated page frame
 - Disk interrupt
 - Page table update
 - Resume process
 - TLB miss
 - Page table fetch

Demand Paging (Before)

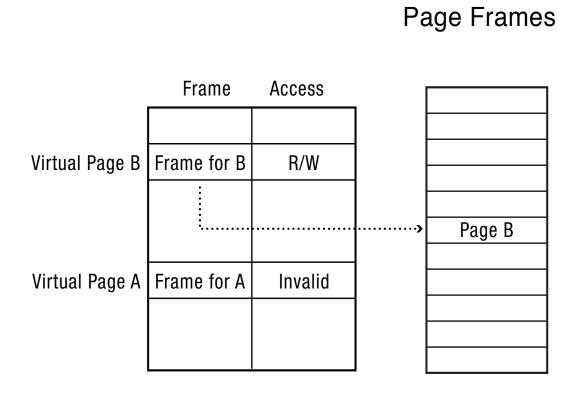


Disk

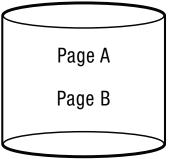


Demand Paging (After)

Page Table

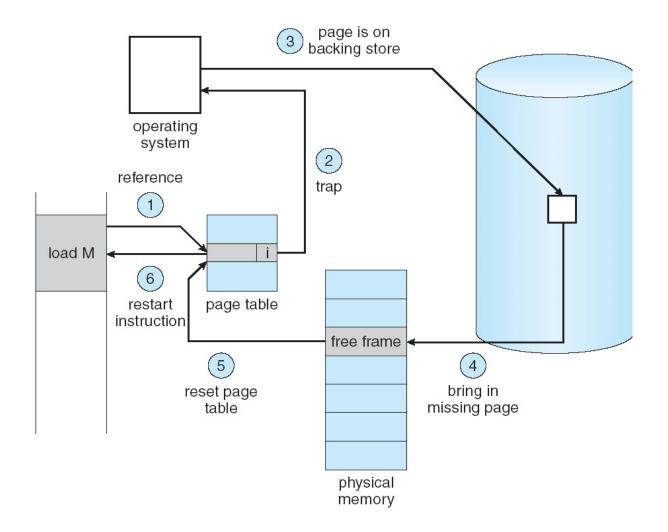


Disk



Physical Memory

Steps In Handling A Page Fault



Allocating a Page Frame

- Select old page to evict
- Find all page table entries that refer to old page
 - If page frame is shared
 - Use a core map
- Set each page table entry to invalid
- Remove any TLB entries
 - Copies of now invalid page table entry
- Write changes on page back to disk, if the evicted page was modified

How Do We Know If Page Has Been Modified?

- Every page table entry has some bookkeeping
 - Has page been modified?
 - Set by hardware on store instruction
 - In both TLB and page table entry
 - Has page been recently used?
 - Set by hardware on in page table entry on every TLB miss
- Bookkeeping info can be reset by the OS kernel
 - When changes to page are flushed to disk
 - To track whether page is recently used

Virtual Memory

From Memory-Mapped Files to Demand-Paged Virtual Memory

- Every process segment backed by a file on disk
 - Code segment -> code portion of executable
 - Data, heap, stack segments -> temp files
 - Shared libraries -> code file and temp data file
 - Memory-mapped files -> memory-mapped files
 - When process ends, delete temp files
- Unified memory management across file buffer and process memory

Virtual Memory Motivation

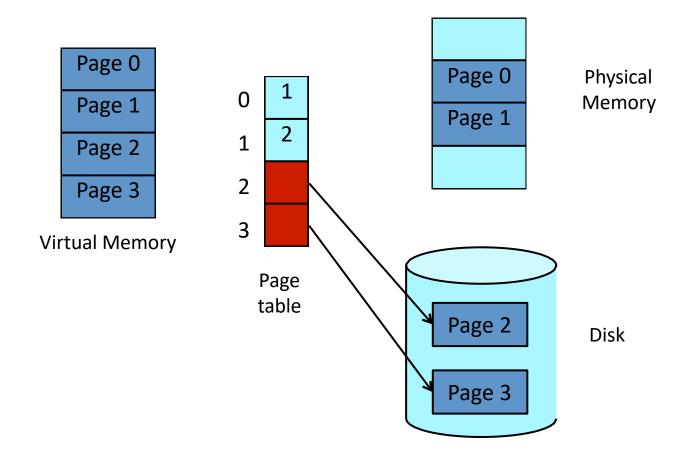
- Previous approach to memory management
 - Must completely load user process in memory
 - One large AS or too many ASes => out of memory
- Observation: locality of reference
 - Temporal: access memory location accessed just now
 - Spatial: access memory location adjacent to locations accessed just now
- Implication: process only needs a small part of address space at any moment!
 - Can load programs faster (don't load everything)
 - Can fit more programs in memory (better utilization)

Virtual Memory Idea

OS and hardware produce illusion of disk as fast as main memory, or main memory as large as disk

- Process runs when not all pages are loaded in memory
 - Only keep referenced pages in main memory
 - Keep unreferenced pages on slower, cheaper backing store (disk)
 - Bring pages from disk to memory when necessary

Virtual Memory Illustration



Virtual Memory Operations

- Detect reference to page on disk
- Recognize disk location of page
- Choose free physical page
 - OS decision: if no free page is available, must replace a physical page
- Bring page from disk into memory
 - OS decision: when to bring page into memory?
- Above steps need hardware and software cooperation

Detect Reference to Page on Disk and Recognize Disk Location of Page

- Overload the present bit of page table entries
- If a page is on disk, clear present bit in corresponding page table entry and store disk location using remaining bits
- Page fault: if bit is cleared then referencing resulting in a trap into OS
- In OS page fault handler, check page table entry to detect if page fault is caused by reference to true invalid page or page on disk

Performance of Demand Paging

- \square Page Fault Rate $0 \le p \le 1$
 - if p = 0 no page faults
 - if p = 1, every reference is a fault
- ☐ Effective Access Time (EAT)

$$EAT = (1 - p) x memory access$$

- + p (page fault overhead
 - + swap page out
 - + swap page in
 - + restart overhead)

Demand Paging Example

- Disparity in memory and disk access times is huge. E.g.,
 - Memory access time = 200 nanoseconds
 - Average page-fault service time = 8 milliseconds
- \blacksquare EAT = $(1 p) \times 200 + p$ (8 milliseconds)

$$= (1 - p) \times 200 + p \times 8,000,000$$

$$= 200 + p \times 7,999,800$$

- ☐ If one out of 1,000 accesses faults, then EAT = 8.2 us, or 40x slower!
- ☐ If want performance degradation < 10 percent
 - □ 200 + 7,999,800 x p < 220, or 7,999,800 x p < 20
 - □ p < .0000025
 - Less than one page fault in every 400,000 memory accesses

OS Decisions

- Page selection
 - When to bring pages from disk to memory?
- Page replacement
 - When no free pages available, must select victim page in memory and throw it out to disk

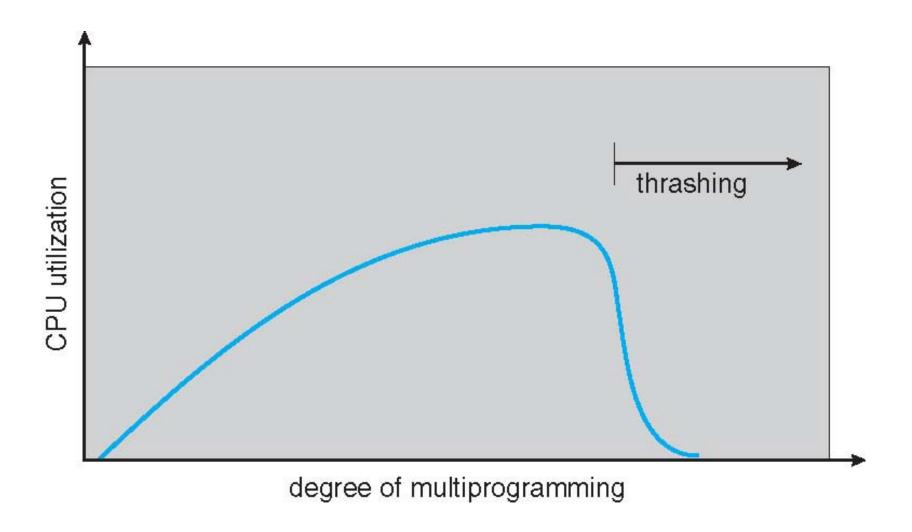
Page Selection Algorithms

- Demand paging: load page on page fault
 - Start up process with no pages loaded
 - Wait until a page absolutely must be in memory
- Request paging: user specifies which pages are needed
 - Requires users to manage memory by hand
 - Users do not always know best
 - OS trusts users (e.g., one user can use up all memory)
- Prepaging: load page before it is referenced
 - When one page is referenced, bring in next one
 - Do not work well for all workloads
 - Difficult to predict future

Thrashing

- What if we need more pages regularly than we have?
 - Page fault to get page
 - □ Replace existing frame
 - But quickly need replaced frame back
- Leads to:
 - High page fault rate
 - Lots of I/O wait
 - Low CPU utilization
 - No useful work done
- \Box Thrashing \equiv system busy just swapping pages in and out

Effects of Thrashing



Page Replacement Implementing LRU (Take 1): Hardware

- A counter for each page
- Every time page is referenced, save system clock (time) into the counter of the page
- Page replacement: scan through pages to find the one with the oldest clock

Problem: have to search all pages/counters!

Implementing LRU (Take 2): Software

- A doubly linked list of pages
- Every time page is referenced, move it to the front of the list
- Page replacement: remove the page from back of list
 - Avoid scanning of all pages
- Problem: too expensive
 - Requires 6 pointer updates for each page reference
 - High contention on multiprocessor

LRU Concept vs. Reality

- □ LRU is considered to be a reasonably good algorithm
- Problem is in implementing it efficiently
 - □ Hardware implementation: counter per page, copied per memory reference, have to search pages on page replacement to find oldest
 - Software implementation: no search, but pointer swap on each memory reference, high contention
- In practice, settle for efficient approximate LRU
 - Find an old page, but not necessarily the oldest
 - LRU is approximation anyway, so approximate more

Clock (Second-Chance) Algorithm

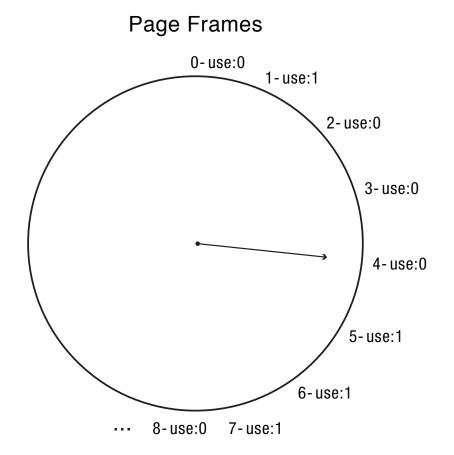
- Goal: remove a page that has not been referenced recently
 - good LRU approximate algorithm
- Idea
 - □ A reference bit per page
 - Memory reference: hardware sets bit to 1
 - Page replacement: OS finds a page with reference bit cleared
 - OS traverses all pages, clearing bits over time

Clock Algorithm Implementation

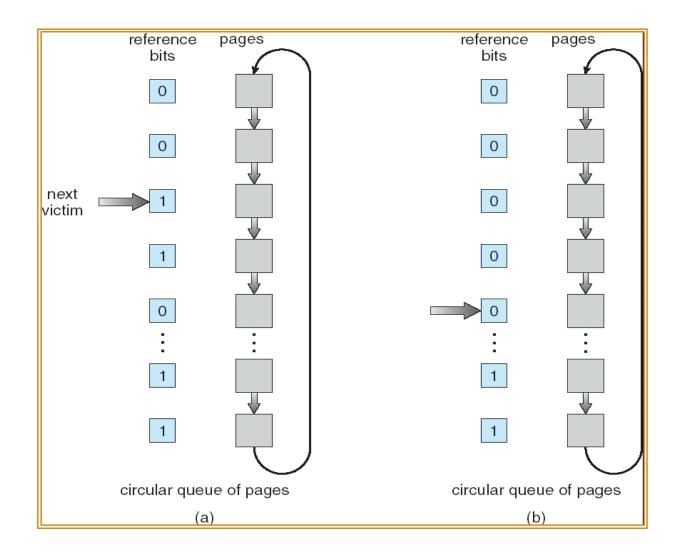
- Combining FIFO with LRU: give the victim page that FIFO selects a second chance
- Keep pages in a circular list = clock
- Pointer to next victim = clock hand
- To replace a page, OS examines the page pointed to by hand
 - □ If ref bit == 1, clear, advance hand
 - Else return current page as victim

Clock Algorithm (aka Second Chance Algorithm): Approximating LRU

- Periodically, sweep through all pages
- If page is unused, reclaim
- If page is used, mark as unused

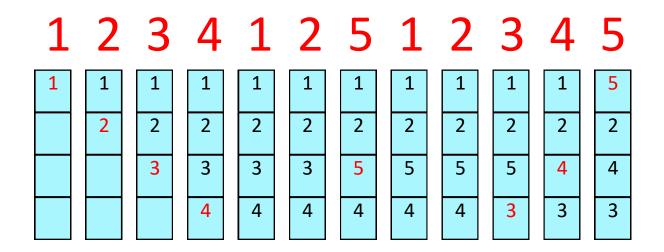


A Single Step in Clock Algorithm



Least Recently Used (LRU) Algorithm


□ Throw out page that hasn't been used in longest time.
Can use FIFO to break ties



8 page faults

Advantage: with locality, LRU approximates Optimal

Clock Algorithm Example



10 page faults

Advantage: simple to implement!

Clock Algorithm Extension

 Problem of clock algorithm: does not differentiate dirty v.s. clean pages

- Dirty page: pages that have been modified and need to be written back to disk
 - More expensive to replace dirty than clean pages
 - One extra disk write (about a few ms)

Clock Algorithm Extension (Cont.)

- Use dirty bit to give preference to dirty pages
- On page reference
 - Read: hardware sets reference bit
 - Write: hardware sets dirty bit
- Page replacement
 - □ reference = 0, dirty = 0 => victim page
 - \square reference = 0, dirty = 1 => skip (don't change)
 - □ reference = 1, dirty = 0 => reference = 0, dirty = 0
 - □ reference = 1, dirty = 1 => reference = 0, dirty = 1
 - advance hand, repeat
 - If no victim page found, run swap daemon to flush unreferenced dirty pages to the disk, repeat