

File Systems (1)



May 28, 2018
Byung-Gon Chun

Acknowledgments. Slides and/or picture in the following are adapted from UW, Columbia, and UC Berkeley slides

So far

- ❑ Operating system kernel - user space/kernel space
- ❑ Computation abstraction – process, thread, synchronization
- ❑ Memory abstraction – address translation, caching, virtual memory

File Systems

- File systems
 - FFS, EXT_x, LFS
- Linux VFS
- RAID, Flash, ...

Outline

- ❑ File system concepts
 - ❑ What is a file?
 - ❑ What operations can be performed on files?
 - ❑ What is a directory and how is it organized?
- ❑ File implementation
 - ❑ How to allocate disk space to files?

What Is A File?

- ❑ User view
 - ❑ Named byte array
 - ❑ Permanently and conveniently available

- ❑ OS view
 - ❑ Map bytes as collection of blocks on physical storage
 - ▶ Stored on nonvolatile storage device: Magnetic Disks, SSDs
 - ▶ Persistent across reboots and power failures

Role of File System

- ❑ Naming
 - ❑ How to “name” files
 - ❑ Translate “name” + offset -> logical block #
- ❑ Reliability
 - ❑ Must not lose file data\
- ❑ Protection
 - ❑ Must mediate file access from different users
- ❑ Disk management
 - ❑ Fair, efficient use of disk space
 - ❑ Fast access to files

File Metadata

- ❑ Metadata: additional system information associated with each file
 - ❑ Name – only information kept in human-readable form
 - ❑ Type of file
 - ❑ Location – pointer to file location on device
 - ❑ Identifier – unique tag (number) identifies file within file system (inode number in UNIX)
 - ❑ Size – current file size
 - ❑ Time – creating, access, modification
 - ❑ Protection – controls who can do reading, writing, executing
 - ❑ Owner and group id
 - ❑ Special file? (directory? Symbolic link?)
- ❑ Meta-data is stored on disk

UNIX File Operations

- ❑ `int creat(const char* pathname, mode_t mode)`
- ❑ `int unlink(const char* pathname)`
- ❑ `int rename(const char* oldpath, const char* newpath)`
- ❑ `int open(const char* pathname, int flags, mode_t mode)`
- ❑ `int read(int fd, void* buf, size_t count);`
- ❑ `int write(int fd, const void* buf, size_t count)`
- ❑ `int lseek(int fd, offset_t offset, int whence)`
- ❑ `int truncate(const char* pathname, offset_t len)`
- ❑ ...

Everything as a File

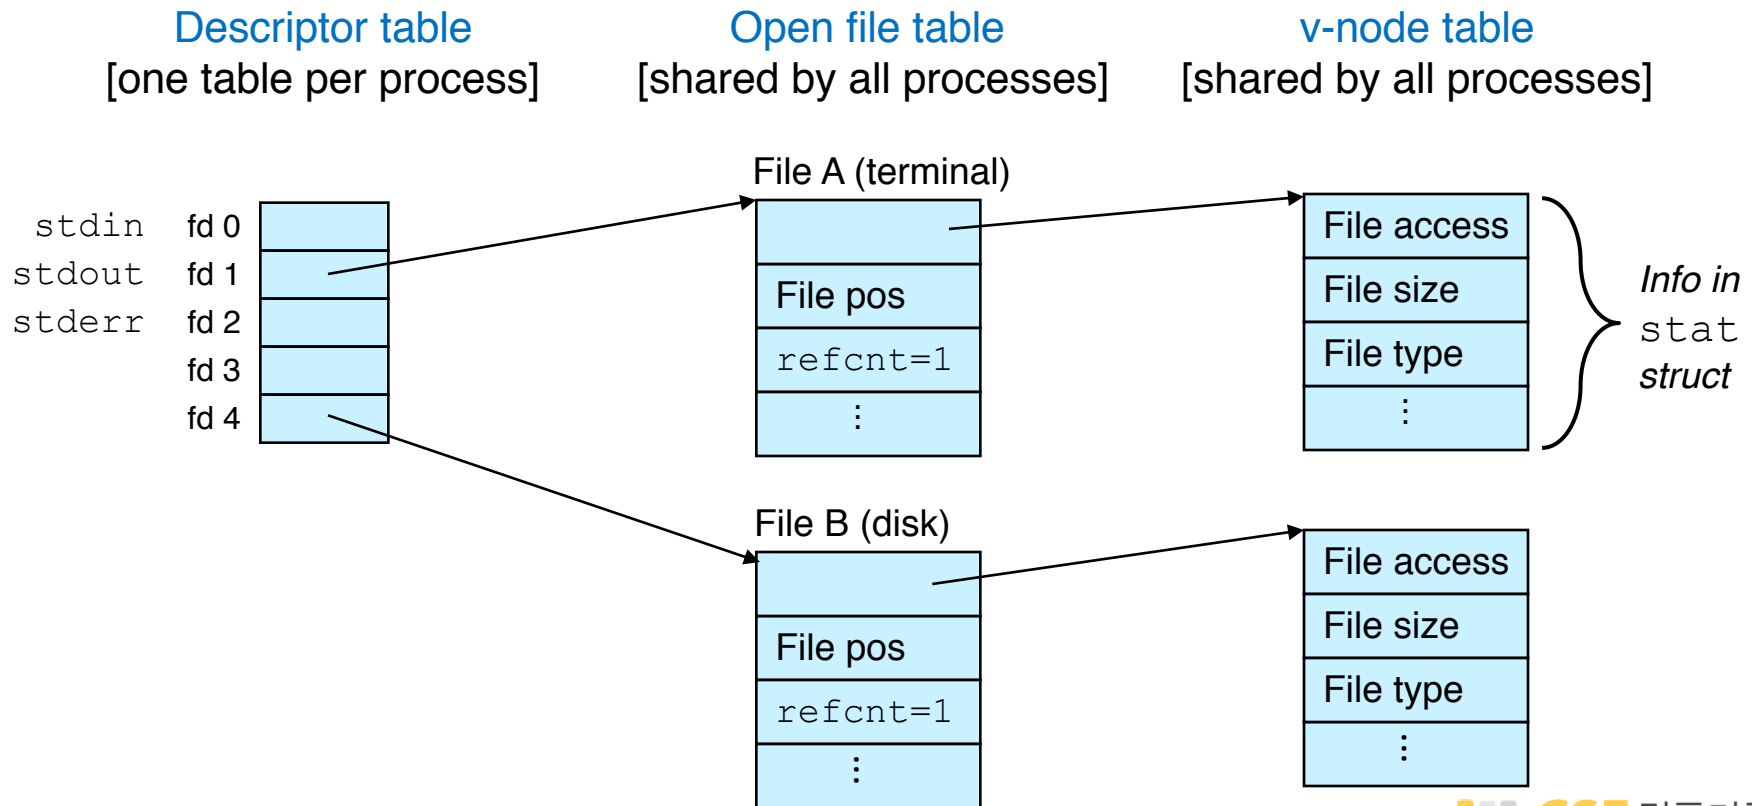
- ❑ A core UNIX tenet from the early days
 - ❑ Block devices (disks, graphics cards in /dev)
 - ❑ Character devices (USB devices, network cards in /dev)
 - ❑ IPC: Pipes, Network sockets
 - ❑ Accessing kernel data structures (/proc, /sys)
 - ❑ Setting kernel configuration
 - ❑ Volatile filesystems in RAM (e.g., tmpfs)
 - ❑ Shared memory (based on tmpfs/shmfs)
 - ❑ Remote files (NFS, SMB, AFP, ...)
 - ❑ Even normal local files
- ❑ Implications
 - ❑ Everything accessed using common API (open, read, write)
 - ❑ Implementation may be totally different
 - ❑ OS must support some measure of **object orientedness**

Open Files

- ❑ Problem: **expensive** to resolve name to identifier on each access
- ❑ Solution: open file before access
 - ❑ **Name resolution**: search directories for file name and check permission
 - ❑ Read relevant file metadata into **open file table** in memory
 - ❑ Return index in open file table (**file descriptor**)
 - ❑ Application pass index to OS for subsequent access
- ❑ **System-wide open file table** shared across processes
- ❑ **Per-process file table** stores current pointer position and index to system-wide open file table

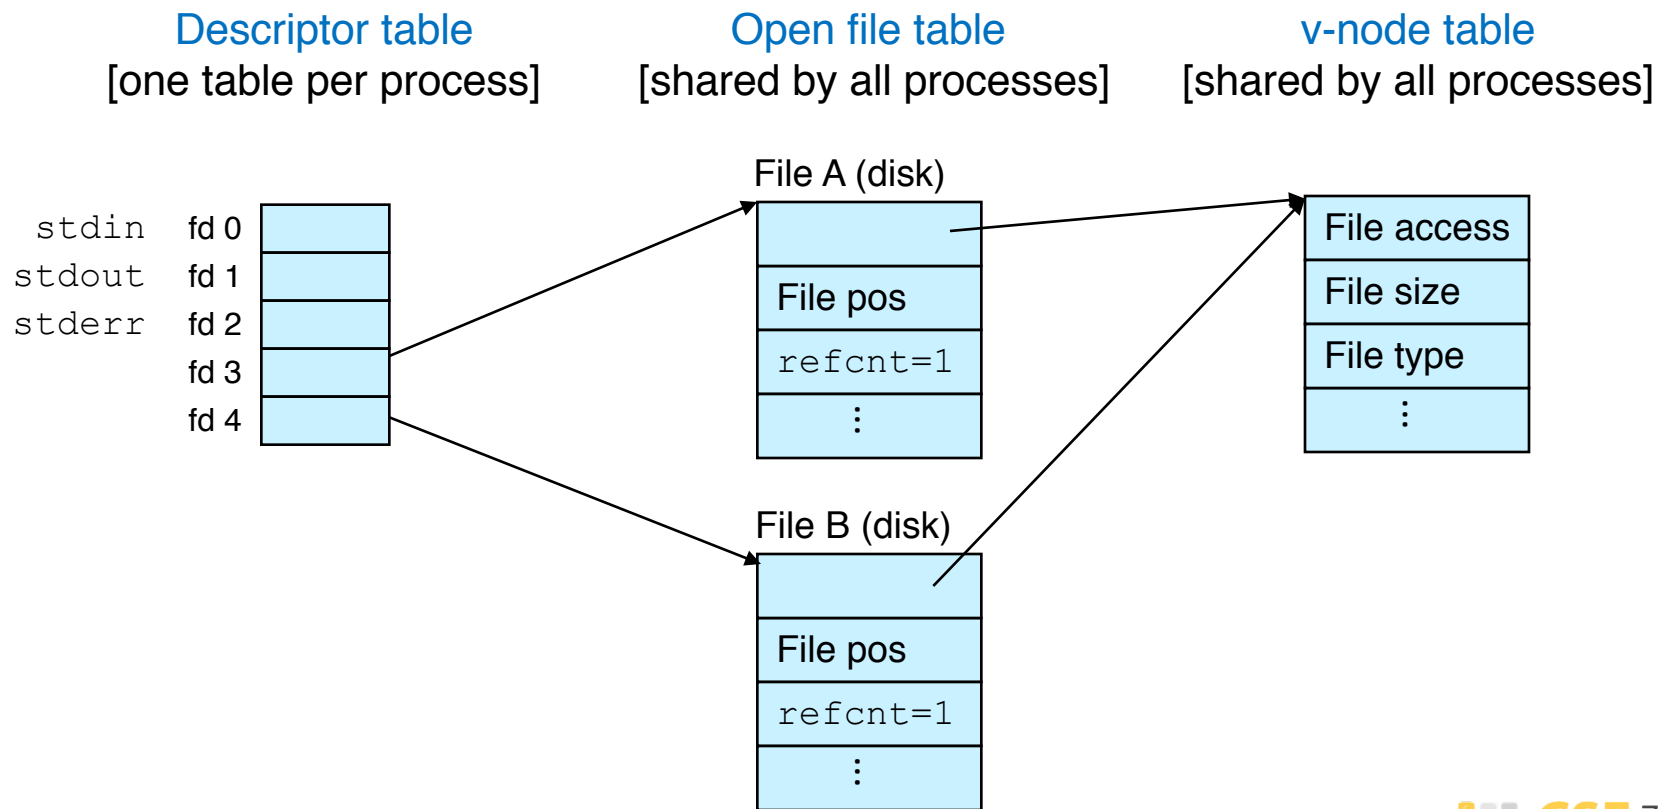
How the Unix Kernel Represents Open Files

- Two descriptors referencing two distinct open disk files. Descriptor 1 (stdout) points to terminal, and descriptor 4 points to open disk file.
- Vnodes: An Architecture for Multiple File System Types in Sun Unix*



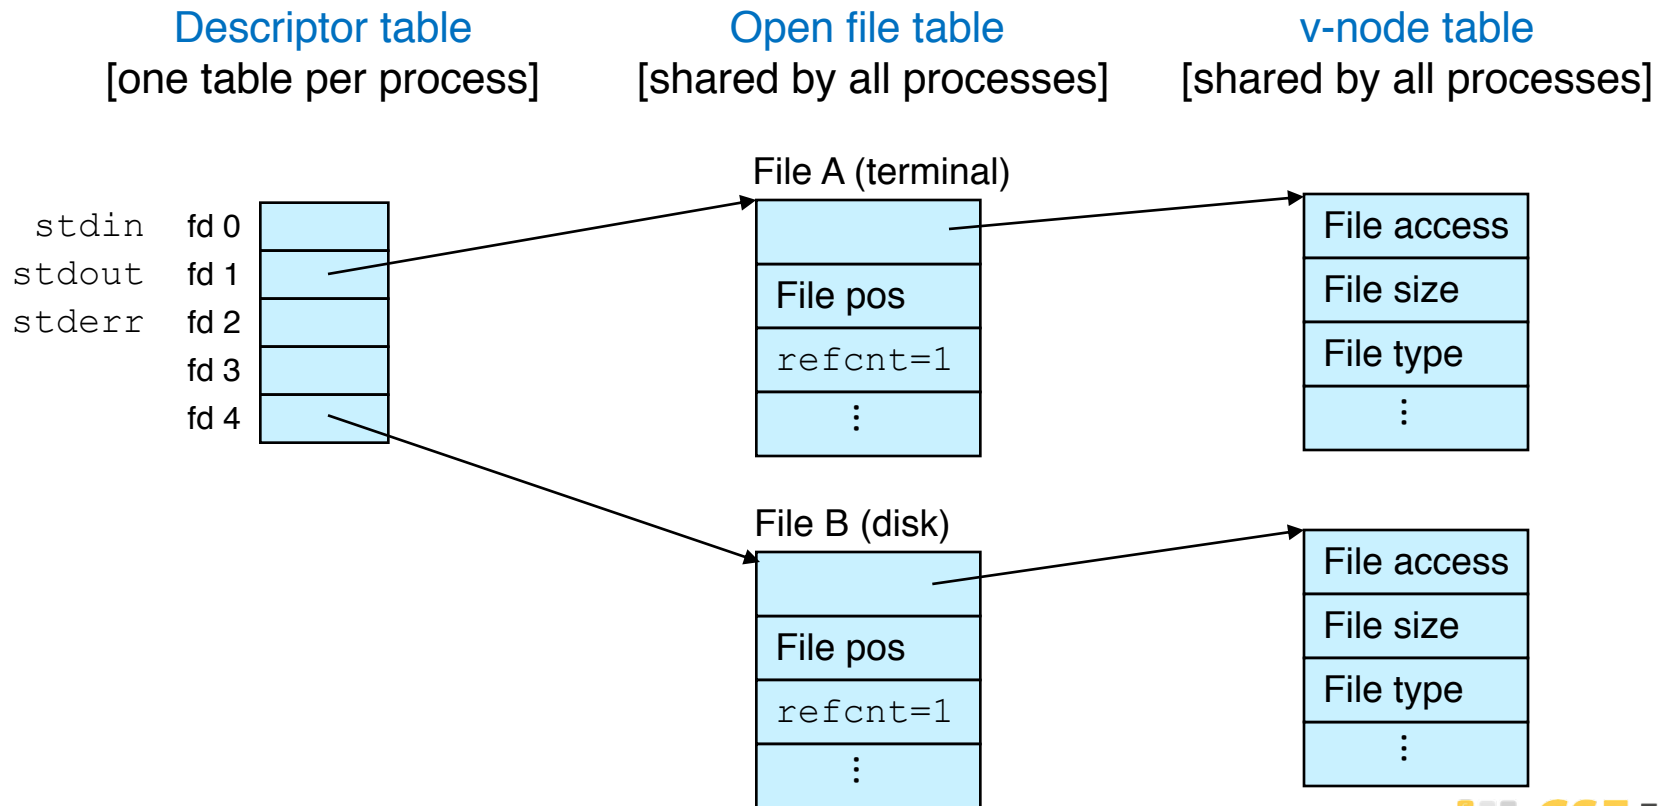
File Sharing

- Two distinct descriptors sharing the same disk file through two distinct open file table entries
 - E.g., Calling open twice with the same filename argument



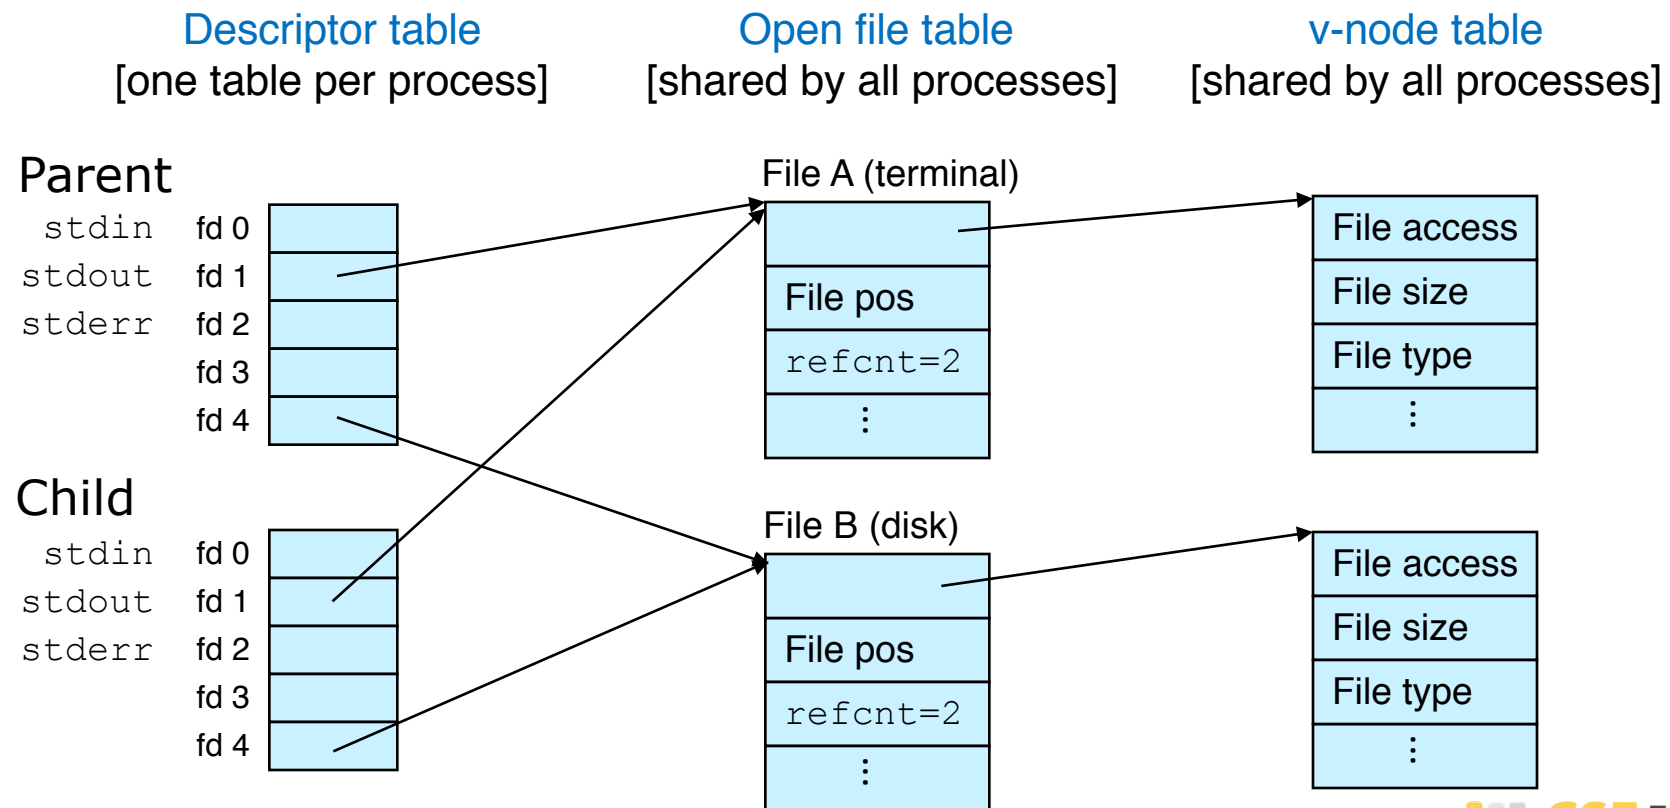
How Processes Share Files: Fork()

- ❑ A child process inherits its parent's open files
- ❑ Before fork() call:



How Processes Share Files: Fork()

- ❑ A child process inherits its parent's open files
- ❑ After fork() call:



Directories

- ❑ Organization technique
 - ❑ Map file name to location on disk
 - ❑ Also stored on disk
- ❑ Single-Level directory
 - ❑ Single directory for entire disk
 - ▶ Each file must have unique name
 - ❑ Not very usable
 - ❑ Special part of disk holds directory listing
- ❑ Two-level directory
 - ❑ Directory for each user
 - ❑ Still not very usable

Tree-Structured Directory

- ❑ Directory stored on disk just like files
 - ❑ Data consists of <name, index> pairs
 - ▶ Index points to file identifier (inode)
 - ▶ Name can be another directory
 - ❑ Designated by special bit in meta-data
 - ❑ Reference by separating names with slashes
 - ❑ Operations
 - ▶ User programs can read (readdir())
 - ▶ Only special system calls can write
- ❑ Special directories
 - ❑ Root (/): fixed index for metadata
 - ❑ . : this directory
 - ❑ .. : parent directory

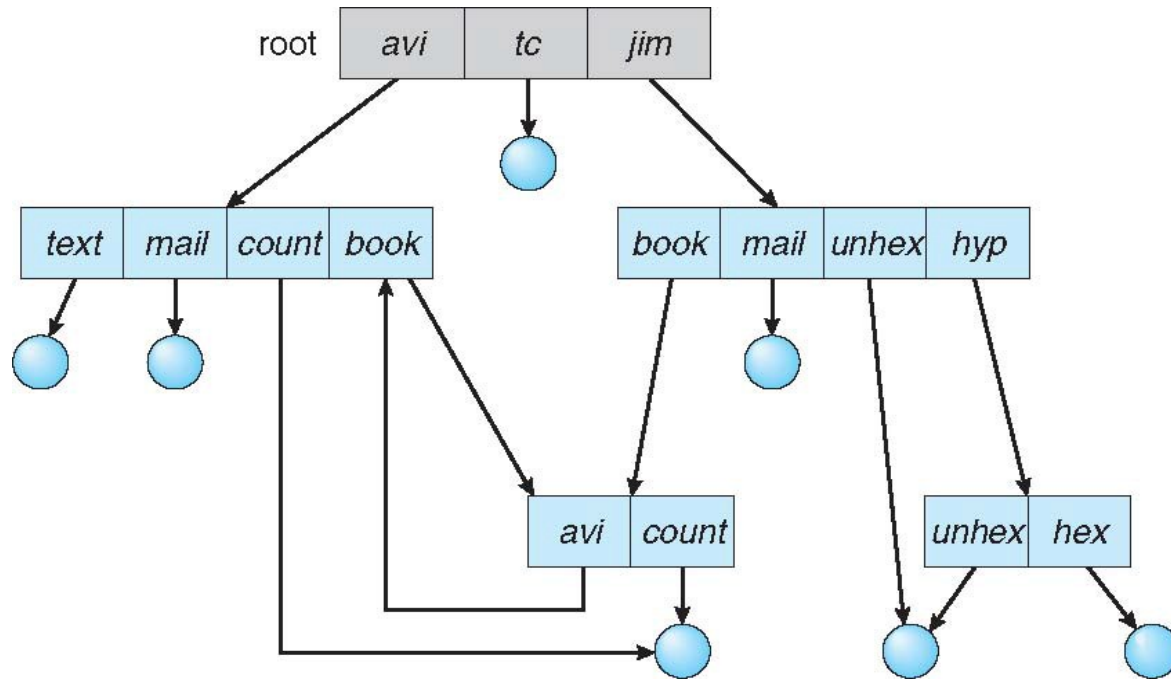
Tree-Structured Directory

- ❑ Example: `mkdir /a/b/c`
 - ❑ Read meta-data 2, look for “a”: find <“a”, 5>
 - ❑ Read 5, look for “b”: find <“b”, 9>
 - ❑ Read 9, verify no “c” exists; allocate c and add “c” to directory

Acyclic-Graph Directories

- ❑ More general than tree structure
 - ❑ Add connections across the tree (no cycles)
 - ❑ Create links from one file (or directory) to another
- ❑ Two types of links
 - ❑ Symbolic link
 - ▶ Special file, designated by bit in meta-data
 - ▶ File data is name to another file
 - ❑ Hard link
 - ▶ Multiple directory entries point to same file
 - ▶ All hard-links are equal: no primary
 - ▶ Store reference count in file metadata
 - ▶ Cannot refer to directories; why?

General Graph Directory and Cycles



- ❑ Cycles cause problems with reference counts
- ❑ E.g., a cycle that isn't accessible through root
- ❑ Need garbage collection

Acyclic-Graph Directories

- ❑ Hard link: “ln a b” (“a” must exist already)
 - ❑ Idea: Can use name “a” or “b” to get to same file data
 - ❑ Implementation: Multiple directory entries point to same meta-data
 - ❑ What happens when you remove a? Does b still exist?
 - ▶ How is this feature implemented???
 - ❑ Unix: Does not create hard links to directories. Why?

Acyclic-Graph Directories

- ❑ Symbolic (soft) link: “ln -s a b”
 - ❑ Can use name “a” or “b” to get to same file data, if “a” exists
 - ❑ When reference “b”, lookup soft link pathname
 - ❑ b: Special file (designated by bit in meta-data)
 - ▶ Contents of b contain name of “a”

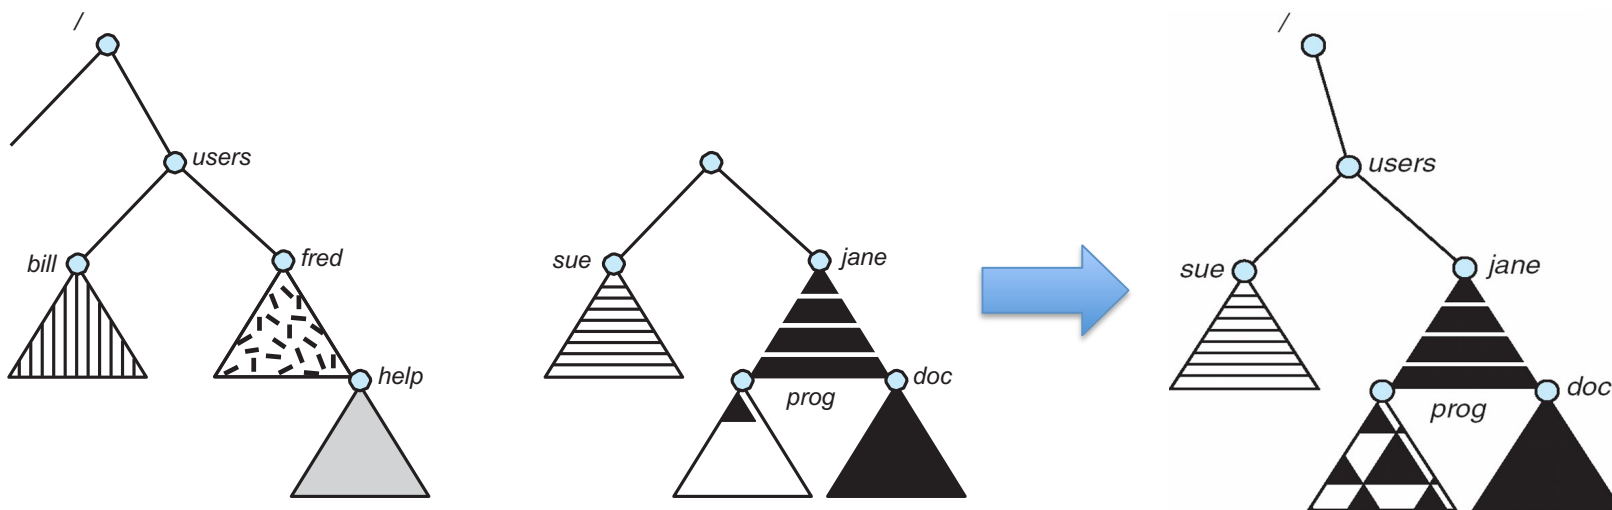
Path Names

- ❑ Absolute path name (full path name)
 - ❑ Start at root directory
 - ▶ E.g. /home/html

- ❑ Relative path name
 - ❑ Full path is lengthy and inflexible
 - ❑ Give each process current working directory
 - ❑ Assume file in current directory

File System Mounting

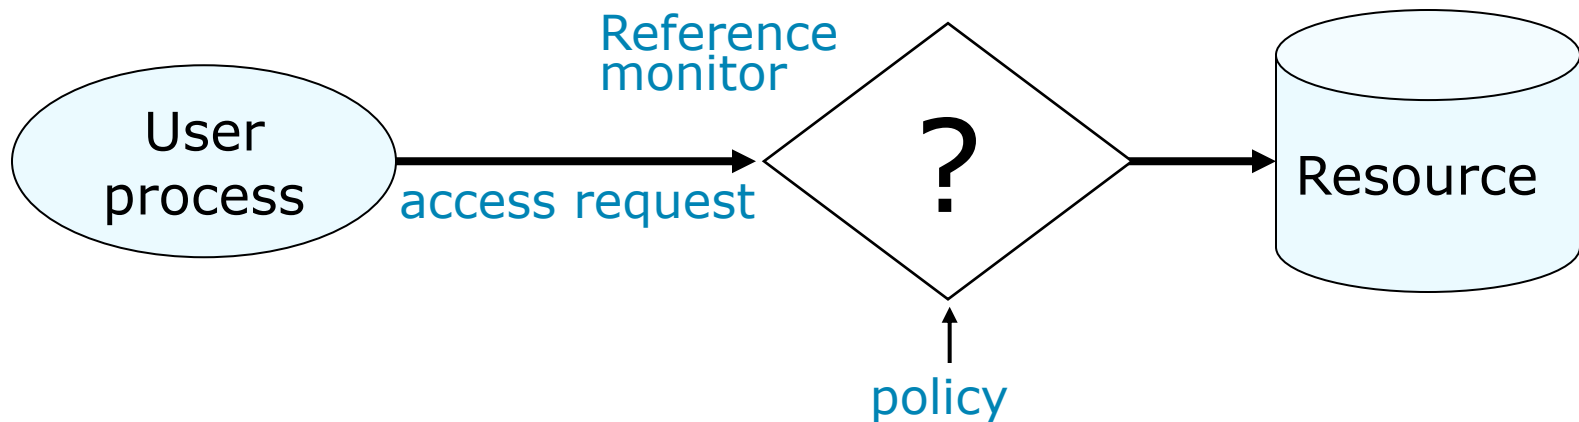
- ❑ Start off with root filesystem
- ❑ New file systems can be mounted into an existing directory (mount point)
- ❑ E.g., `mount -o opts -t ext2 /dev/hda3 /users`



Access control

□ Assumptions

- System knows who the user is
 - ▶ Authentication via name and password, other credential
- Access requests pass through gatekeeper (reference monitor)
 - ▶ System must not allow monitor to be bypassed



Access Control [Lampson]

- ❑ The guard evaluates a function:
permissions = policy(subject, object)
- ❑ If functions are too mathematical, call it an access matrix
(Lampson 1971)

Access control matrix [Lampson]

Objects

Subjects

	File 1	File 2	File 3	...	File n
User 1	read	write	-	-	read
User 2	write	write	write	-	-
User 3	-	-	-	read	read
...					
User m	read	write	read	write	read

Implementation concepts

- ❑ Access control list (ACL)

- ❑ Store column of matrix with the resource

- ❑ Capability

- ❑ User holds a “ticket” for each resource

- ❑ Two variations

- ▶ store row of matrix with user, under OS control
 - ▶ unforgeable ticket in user space

	File 1	File 2	...
User 1	read	write	-
User 2	write	write	-
User 3	-	-	read
...			
User m	Read	write	write

Access control lists are widely used, often with groups

Some aspects of capability concept are used in many systems

ACL vs Capabilities

- ❑ Access control list
 - ❑ Associate list with each object
 - ❑ Check user/group against list
 - ❑ Relies on authentication: need to know user
- ❑ Capabilities
 - ❑ Capability is unforgeable ticket
 - ▶ Random bit sequence, or managed by OS
 - ▶ Can be passed from one process to another
 - ❑ Reference monitor checks ticket
 - ▶ Does not need to know identify of user/process

Protection

❑ Type of access

- ❑ Read, write, execute, append, delete, list, ...

❑ Access control list

- ❑ Associate lists of users with access rights for every file
- ❑ Advantage: complete control
- ❑ Disadvantage
 - ▶ Tedious to construct list (may not know in advance for all users)
 - Require variable-size information

❑ Classify users

- ❑ Assign a owner and group to each file
- ❑ Different permissions based on who is accessing: owner, group, other
- ❑ Advantage: easier to implement
- ❑ Disadvantage: no fine grained control