

File Systems (2)



May 30, 2018
Byung-Gon Chun

Acknowledgments. Slides and/or picture in the following are adapted from UW, Columbia, and UC Berkeley slides

Outline

☐ File system concepts

- ☐ What is a file?
- ☐ What operations can be performed on files?
- ☐ What is a directory and how is it organized?

☐ File implementation

- ☐ How to allocate disk space to files?

Typical File Access Patterns

❑ Sequential Access

- ❑ Data read or written in order
 - ▶ Most common access pattern
 - E.g., copy files, compiler read and write files,
- ❑ Can be made very fast (peak transfer rate from disk)

❑ Random Access

- ❑ Randomly address any block
 - ▶ E.g., update records in a database file
- ❑ Difficult to make it fast (**seek time and rotational delay**)

Disk Management

- ❑ Need to track where file data is on disk
 - ❑ How should we map logical sector # to surface #, track #, and sector #?
 - ▶ Order disk sectors to minimize seek time for sequential access
- ❑ Need to track where file metadata is on disk
- ❑ Need to track free versus allocated areas of disk
 - ❑ E.g., block allocation bitmap (Unix)
 - ▶ Array of bits, one per block
 - ▶ Usually keep entire bitmap in memory

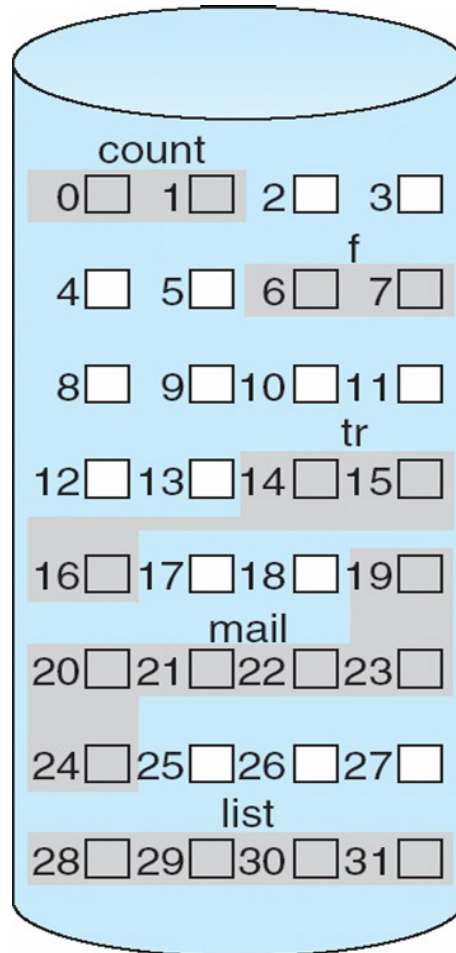
Allocation Strategies

- ❑ Various approaches (similar to memory allocation)
 - ❑ Contiguous
 - ❑ Extent-based
 - ❑ Linked
 - ❑ FAT tables
 - ❑ Indexed
 - ❑ Multi-level indexed
- ❑ **Key metrics**
 - ❑ Fragmentation (internal & external)?
 - ❑ Grow file over time after initial creation ?
 - ❑ Fast to find data for sequential and random access?
 - ❑ Easy to implement?
 - ❑ Storage overhead?

Contiguous Allocation

- ❑ Allocate files like **continuous memory allocation** (base & limit)
 - ❑ User specifies length, file system allocates space all at once
 - ❑ Can find disk space by examining bitmap
 - ❑ Metadata: contains starting location and size of file

Contiguous Allocation Example



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

Pros and Cons

❑ Pros

- ❑ Easy to implement
- ❑ Low storage overhead
(two variables to specify disk area for file)
- ❑ Fast sequential access since data stored in continuous blocks
- ❑ Fast to compute data location for random addresses.
- ❑ Just an array index

❑ Cons

- ❑ Large external fragmentation
- ❑ Difficult to grow file

Extent-based Allocation

- ❑ Multiple contiguous regions per file (like segmentation)
 - ❑ Each region is an **extent**
 - ❑ Metadata: contains small array of entries designating extents
 - ▶ Each entry: start and size of extent

Pros and Cons

❑ Pros


- ❑ Easy to implement
- ❑ Low storage overhead (a few entries to specify file blocks)
- ❑ File can grow overtime (until run out of extents)
- ❑ Fast sequential access
- ❑ Simple to calculate random addresses

❑ Cons


- ❑ Help with external fragmentation, but still a problem

Linked Allocation

- ❑ All blocks (fixed-size) of a file on linked list
 - ❑ Each block has a pointer to next
 - ❑ Metadata: pointer to the first block



Linked Allocation Example



Pros and Cons

❑ Pros

- ❑ No external fragmentation
- ❑ Files can be easily grown with no limit
- ❑ Also easy to implement, though awkward to spare space for disk pointer per block


❑ Cons

- ❑ Large storage overhead (one pointer per block)
- ❑ Potentially slow sequential access
- ❑ Difficult to compute random addresses

Variation: FAT table

- ❑ Store linked-list pointers outside block in File-Allocation Table
 - ❑ One entry for each block
 - ❑ Linked-list of entries for each file
- ❑ Used in MSDOS and Windows operating systems

FAT Example



Pros and Cons

❑ Pros

- ❑ Fast random access. Only search cached FAT


❑ Cons

- ❑ Large storage overhead for FAT table
- ❑ Potentially slow sequential access

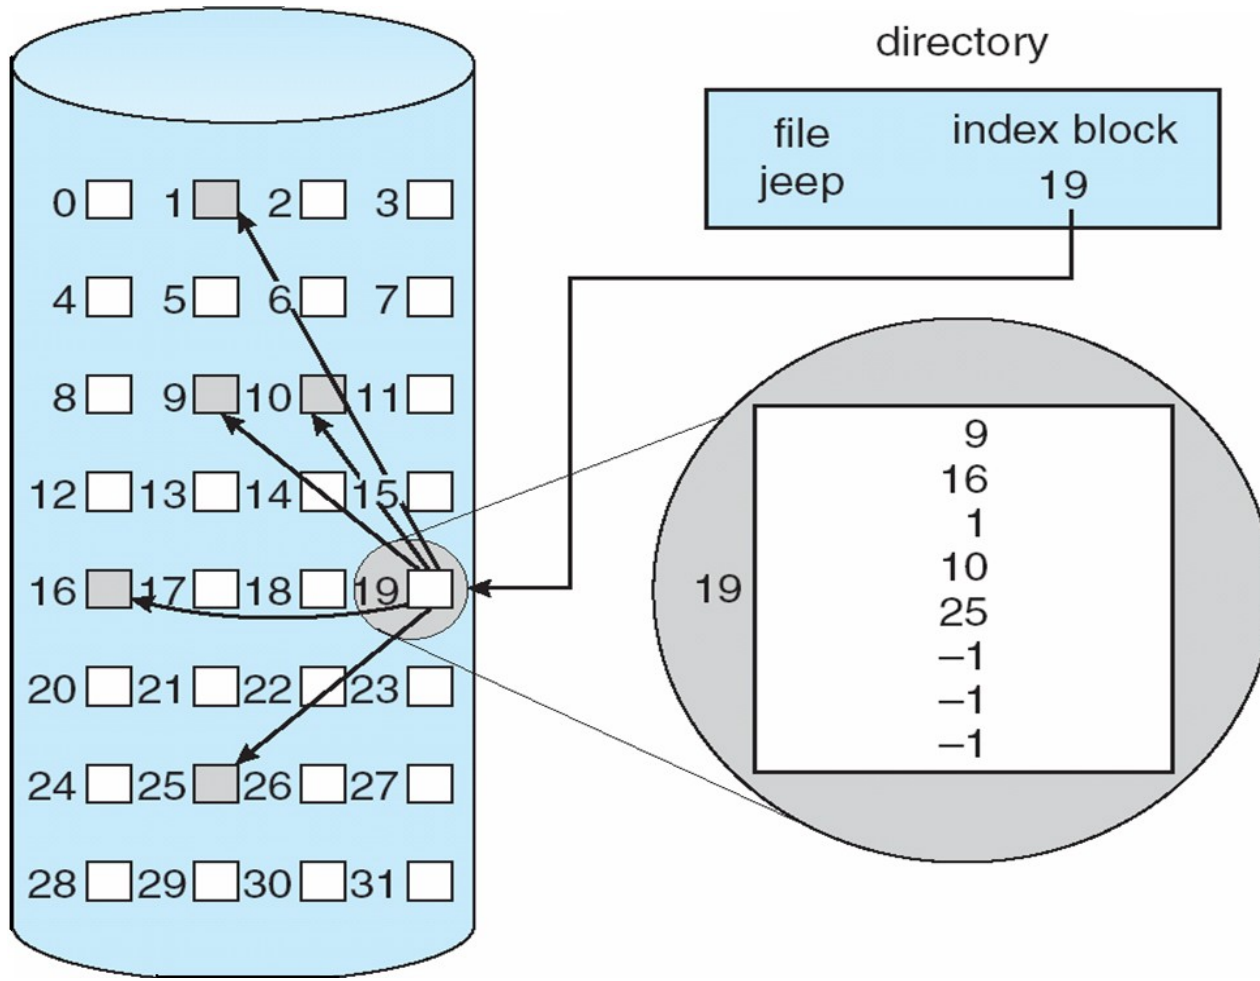
Indexed Allocation

- ❑ File has array of pointers (index) to block
 - ❑ Allocate block pointers contiguously in metadata
 - ▶ Must set max length when file created
 - ▶ Allocate pointers at creation, allocate blocks on demand
 - ▶ Cons: fixed size, same overhead as linked allocation
 - ❑ Maintain multiple lists of block pointers
 - ▶ Last entry points to next block of pointers
 - ▶ Cons: may need to access a large number of pointer blocks

block
pointers




Indexed Allocation Example

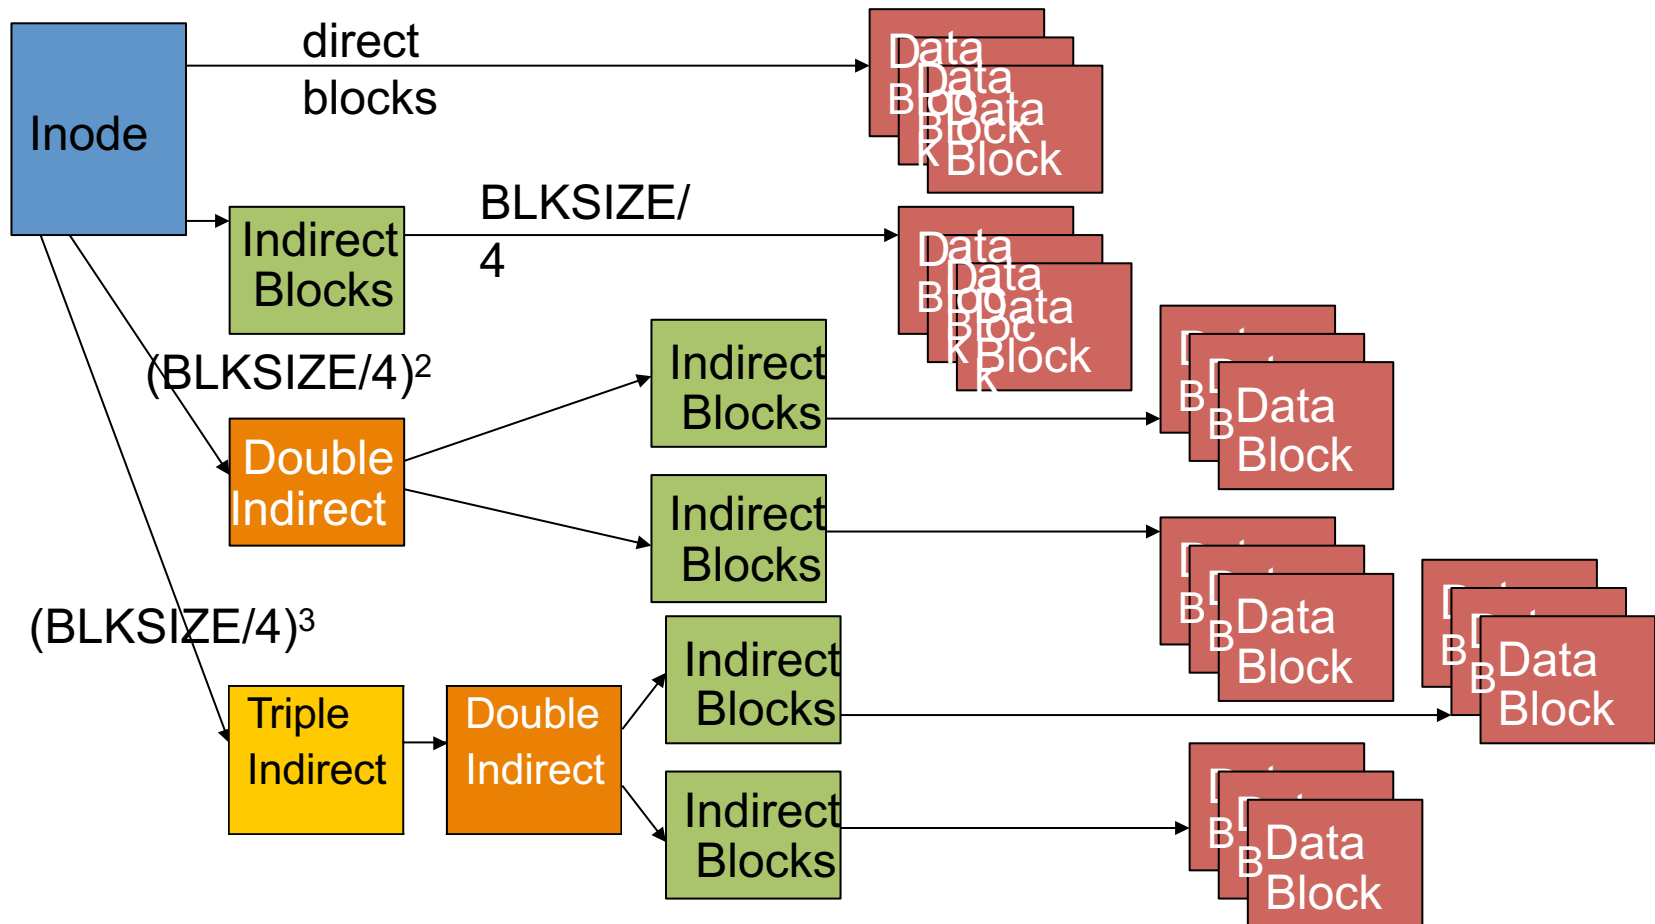


Multi-level Indexed Files

- ❑ Block index has multiple levels




Multi-level Indexed Allocation (UNIX FFS, and Linux ext2/ext3)



Advanced Data Structures


- ❑ More sophisticated data structures
- ❑ B+ Trees
 - ❑ Used by many high perf filesystems for directories and/or data
 - ❑ E.g., XFS, ReiserFS, ext4, MSFT NTFS and ReFS, IBM JFS, brzs
 - ❑ Can support very large files (including sparse files)
 - ❑ Can give very good performance (minimize disk seeks to find block)



Free Space Management

❑ File system maintains **free-space list** to track available blocks/clusters

❑ **Free bitmap** stored in the superblock



$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

❑ Linked free list

- Pros: space efficient
- Cons: requires many disk reads to find free cluster of right size

❑ Grouping

- Use a free index-block containing n-1 pointers to free blocks and a pointer to the next free index-block

❑ Counting

- Free list of variable sized contiguous clusters instead of blocks
- Reduces number of free list entries

