

Project 3

SPL, SNU



Project 3 Overview

- Design and implement WRR (Weighted Round-Robin) scheduler working in your ARTIK
 - Define and Implement a new scheduler
 - Implement load balancing mechanism
 - Make the scheduler class as the default scheduler for init process
 - For both systemd & kthreadd
 - Examine the scheduler performance with Trial
 - Improve the WRR scheduler(for 4 members)
 - Open question
 - Extra.md

WRR Scheduler

Linux Scheduler Basic

- Multi-level scheduling
 - Real-time tasks has priority over other tasks
- Real-time tasks are scheduled in FCFS or RR fashion
- Other tasks are scheduled by CFS (Completely-Fair Scheduler) algorithm
- Each CPU maintains separate run queues for tasks
 - To prevent contention while accessing run queue

WRR Scheduler

- Weighted Round-Robin Scheduler
- Tasks are executed in round-robin fashion, but gets different time slices according to their weights
 - Default weight is 10
 - $\text{Time slice} = \text{Weight} * 10\text{ms}$
- WRR has higher priority than CFS, but lower priority than RT (Real-Time) scheduler
- Load balancing should be implemented

Multi-level Run Queue with WRR

Run Queue per CPU (struct rq)

High
Priority



Low
Priority



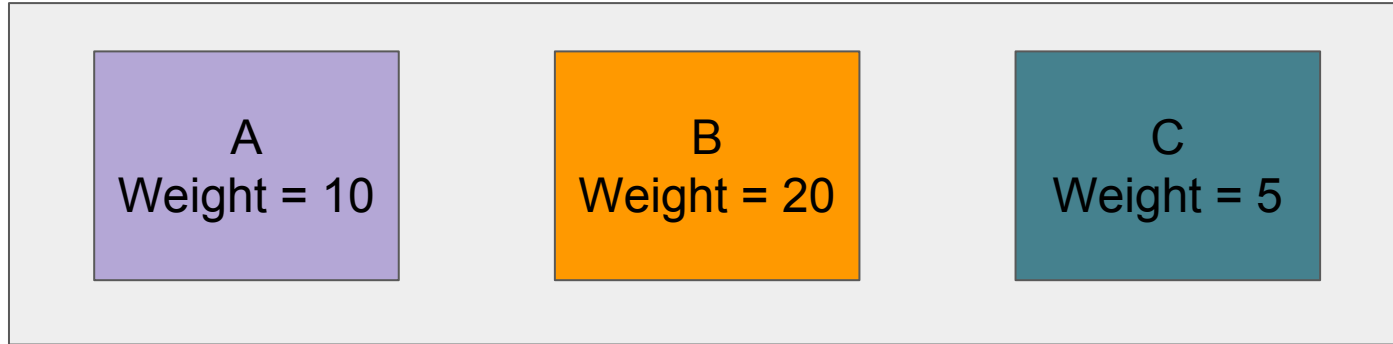
Real-Time Run Queue (struct rt_rq)

WRR Run Queue (struct wrq_rq)

CFS (Completely fair scheduler) Run Queue (struct cfs_rq)

WRR Scheduling Example

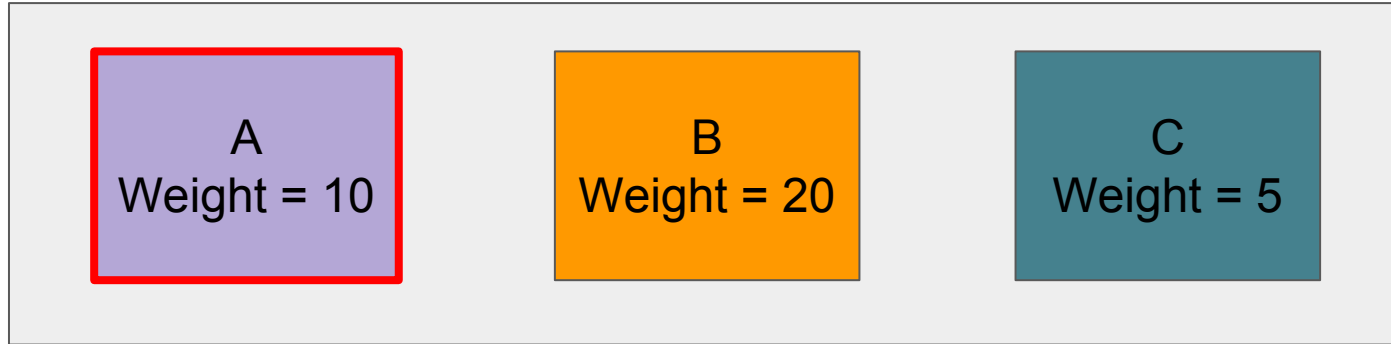
Three tasks are currently in WRR run queue



WRR Scheduling Example

0ms passed

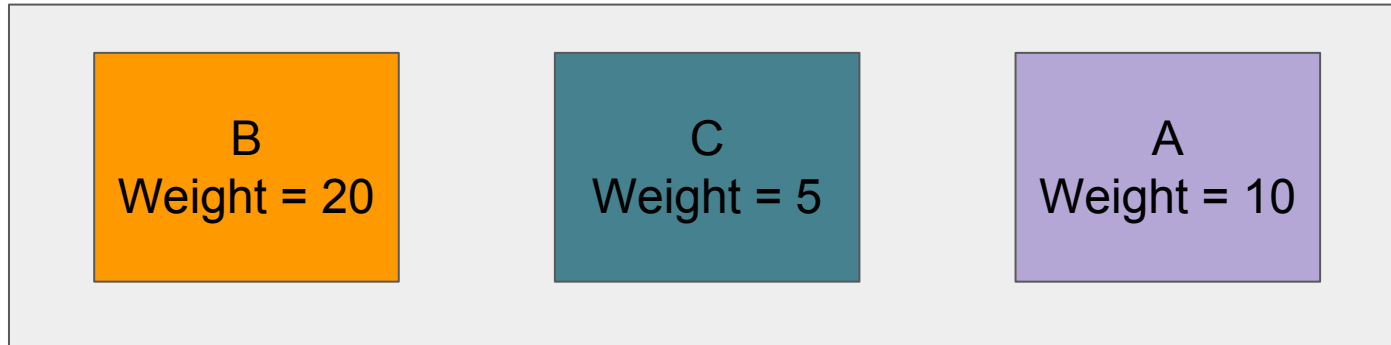
A starts running first



WRR Scheduling Example

100ms passed

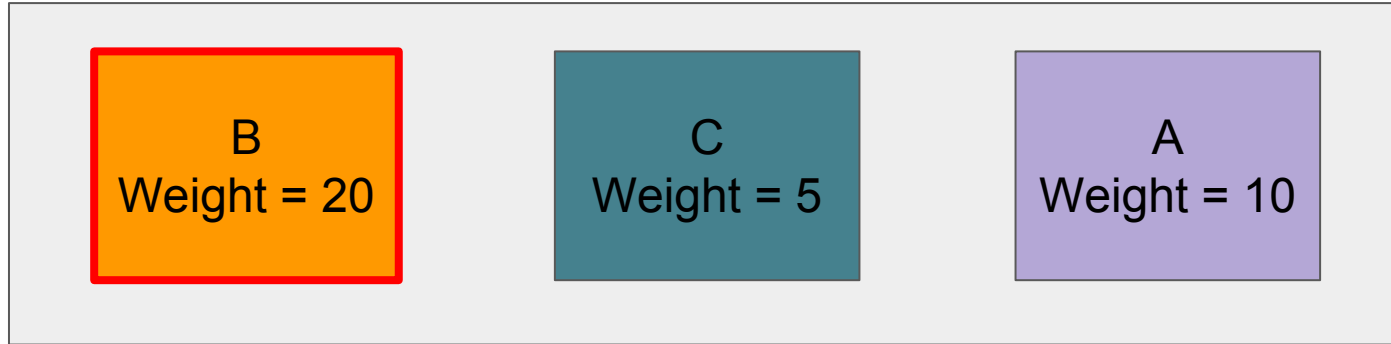
A stopped, and moved to the tail of the run queue because the task is not finished



WRR Scheduling Example

100ms passed


The next task (B) starts running



WRR Scheduling Example

200ms passed

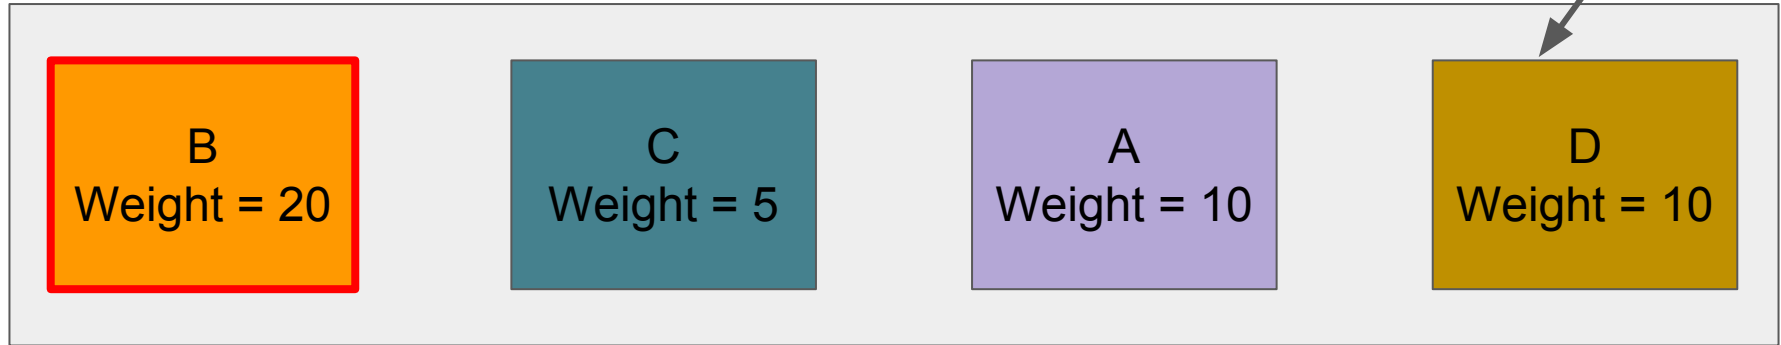
B is still running and not stopped, because it got a 200ms time slice



WRR Scheduling Example

250ms passed

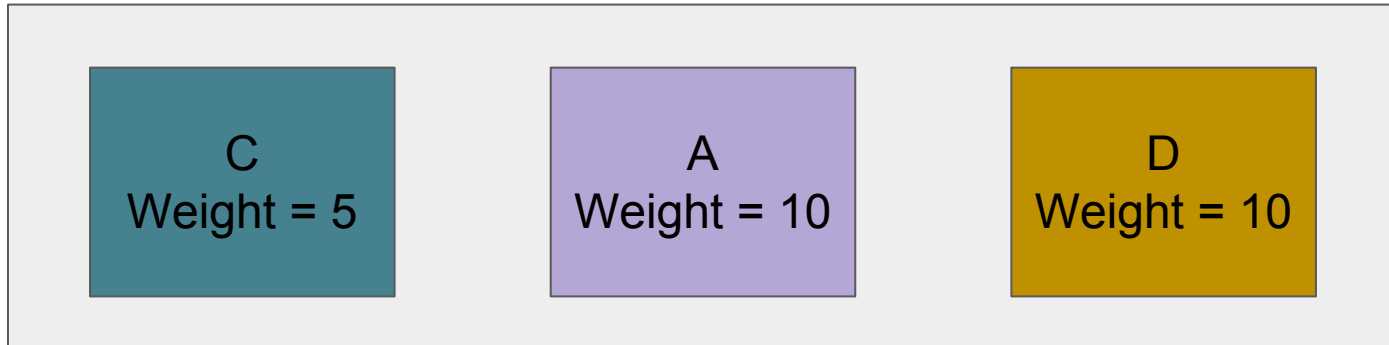
D comes in, and it is added to the tail of the run queue



WRR Scheduling Example

280ms passed

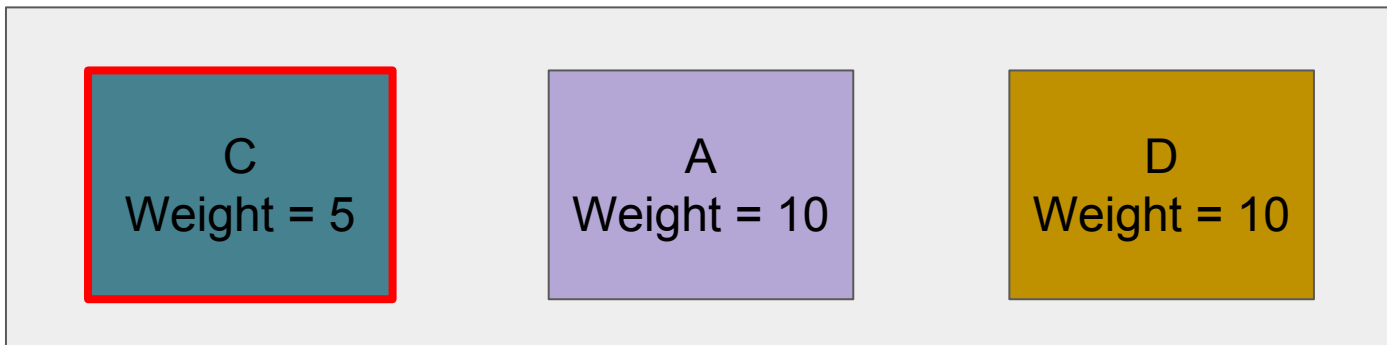
B has been finished, and removed from the run queue because there is no more work left B



WRR Scheduling Example

280ms passed

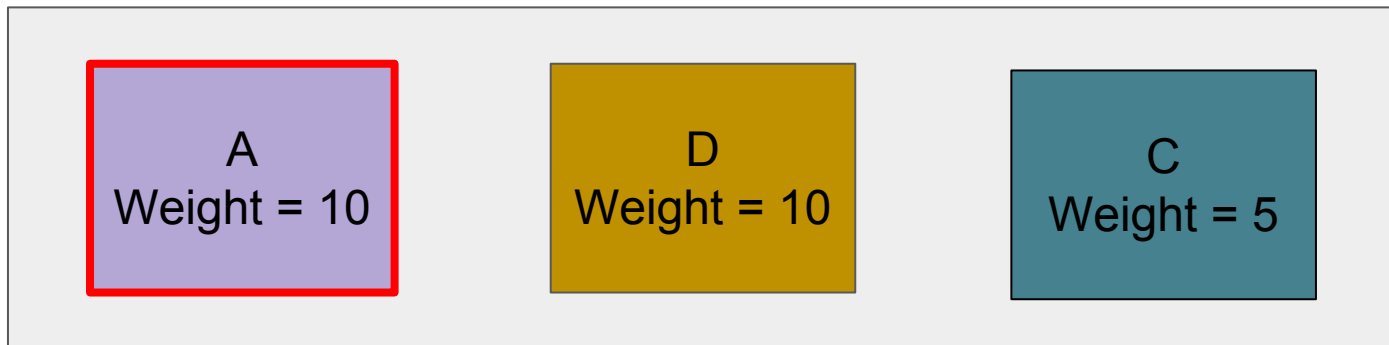
C starts running



WRR Scheduling Example

330ms passed

C has been stopped and moved to the tail. A starts running again



Load balancing


- Balance loads among each CPU's run queue
- Make sure that it only works when more than one CPUs are active
 - When there is no heavy task, only one CPU is active with high probability!
 - CPU hotplug
 - `for_each_online_cpu(cpu)`
- Should be done every 2000ms

Load balancing algorithm

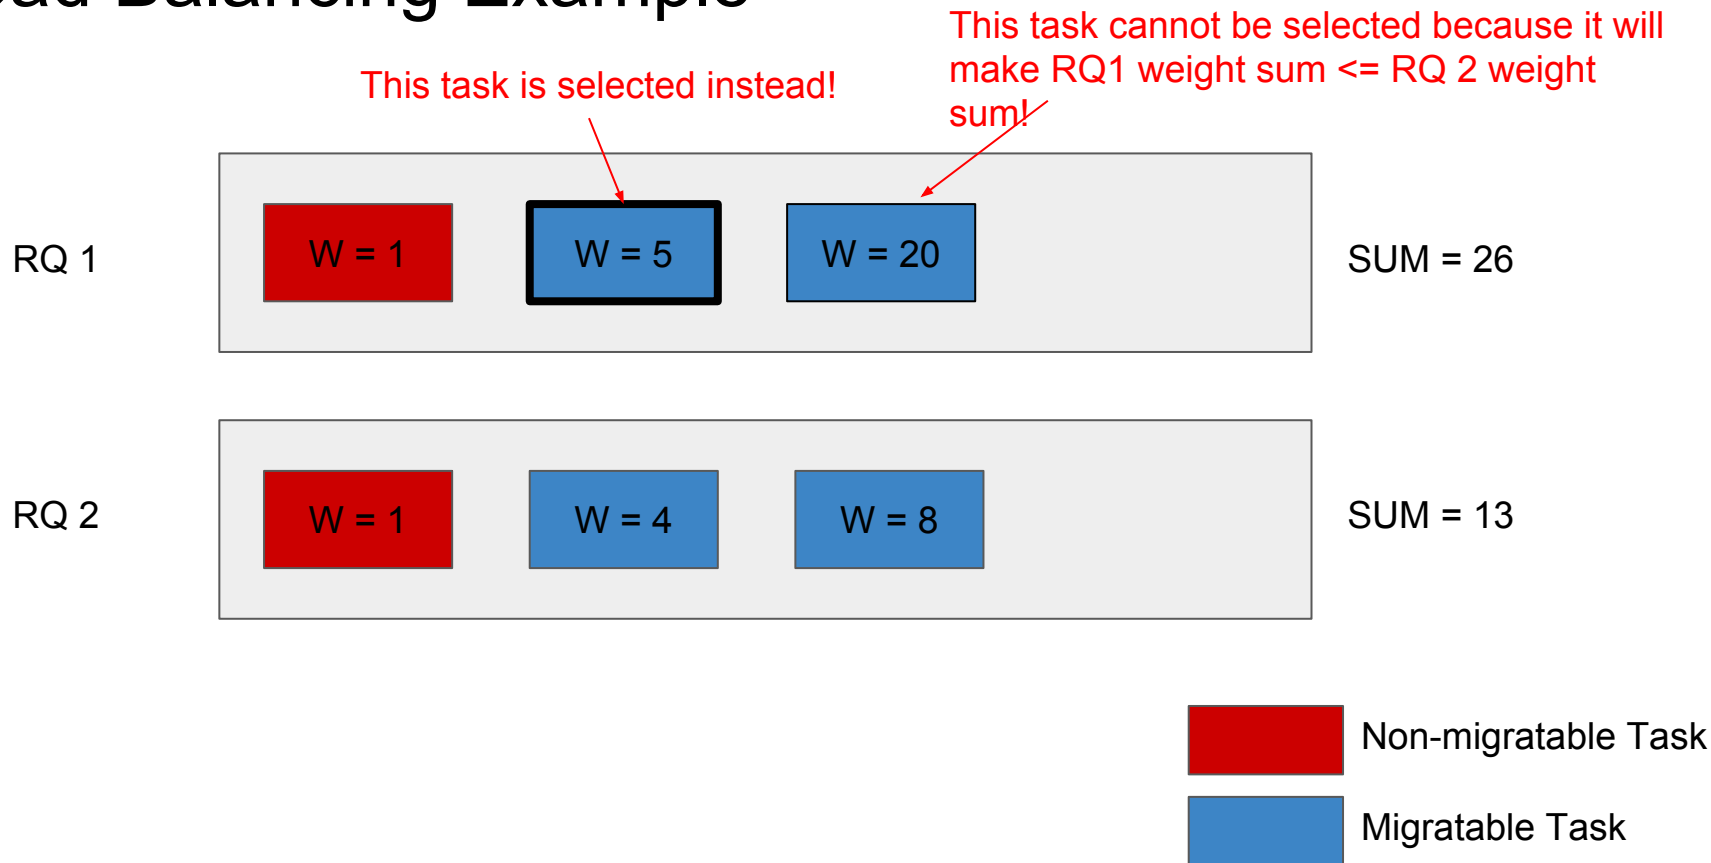
- Pick two run queues with the minimum weight sum and the maximum weight sum, respectively.
 - Let's call them RQ_MIN and RQ_MAX
- Pick a task with the biggest weight among tasks meeting these conditions:
 - The picked task should be able to be migrated to RQ_MIN
 - Migration should not cause weight of RQ_MIN to become **bigger than or equal** to RQ_MAX
 - The currently running task cannot be picked
- Migrate the task if an eligible task exists (there may be no eligible tasks)

Load Balancing Example

Migrate a task from RQ 1 to RQ 2



Load Balancing Example



Load Balancing Example

After migration...

RQ 1

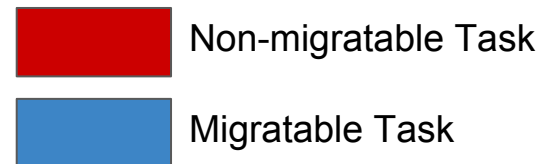


SUM = 21

RQ 2



SUM = 18



Load Balancing Example

Cannot migrate a task
from RQ1 to RQ2!

Migration of this task makes weight sum of
RQ2 \geq weight sum of RQ1

RQ 1

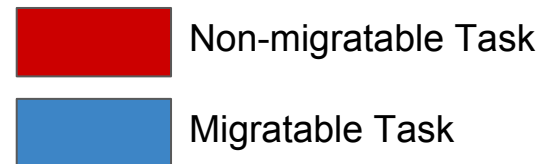


SUM = 21

RQ 2



SUM = 18




Load Balancing Example

Can be migrated now!

Non-migratable task has been changed!
Ex) Context switching

RQ 1

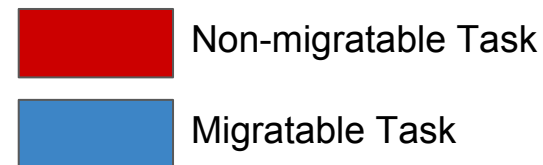


SUM = 21

RQ 2




SUM = 18



Load Balancing Example

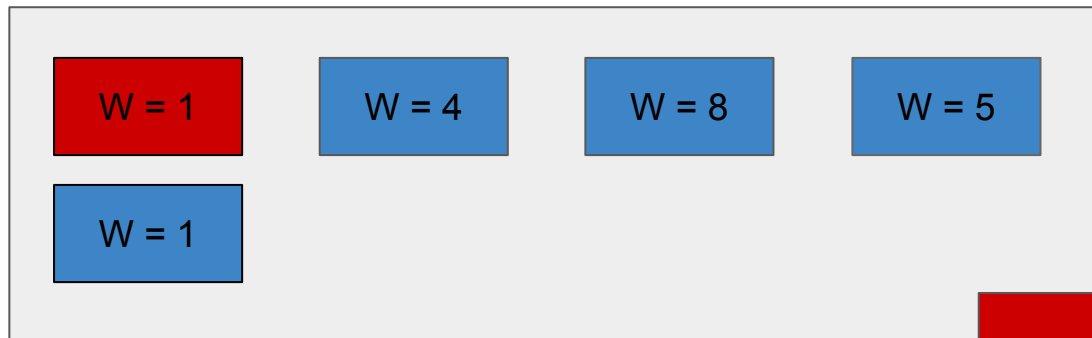
After migration...

RQ 1



SUM = 20

RQ 2



SUM = 19



Non-migratable Task



Migratable Task

Scheduler Implementation

Before Start...

- Modify arch/arm/configs/artik10_defconfig
 - Find CONFIG_SCHED_DEBUG and make it to “CONFIG_SCHED_DEBUG=y”
 - You need this option to debug your scheduler
 - Performance degradation could happen
 - (Optional) Enable CONFIG_SCHEDSTATS for more detailed debugging

Implementation Overview (1)

- Define necessary constants and data structures
 - `include/linux/sched.h`
 - `include/uapi/linux/sched.h`
 - ...
- Register a new scheduler class for WRR and implement necessary functions in **`kernel/sched/wrr.c`**
- Modify **`kernel/sched/debug.c`** to print additional necessary information about WRR scheduler
 - Optionally `kernel/sched/stats.c` too

Implementation Overview (2)

- Modify **kernel/sched/core.c** to support WRR
 - Ex) Trigger load balancing function
 - You might need to register function signatures of `wrr.c` in `kernel/sched/sched.h` for them to be used in other files (ex: `core.c`)
- Implement necessary system calls, `sched_setweight` & `sched_getweight`
- After confirming WRR is working, make WRR as the default scheduling policy of `init` & `kthread`
 - `include/linux/init_task.h`
 - `kernel/kthread.c`

Define constants and data structures

- Define SCHED_WRR as **6**
 - include/**uapi**/linux/sched.h
- Define WRR scheduler information inside struct task_struct
 - Like RT or CFS
 - List head for putting into WRR run queue
 - weight, time slice, ...
- Define run queue for tasks scheduled by WRR
 - “struct rq” should also have WRR run queue
 - struct rq is CPU run queue
 - What information should it have?
 - Should this have locking mechanism inside it?

Register a new scheduler for WRR

- Declare and define **wrr_sched_class** in kernel/sched/sched.h and in kernel/sched/wrr.c
 - Take a look at kernel/sched/fair.c & kernel/sched/rt.c
 - Its next scheduler class should be fair_sched_class because it has higher priority than that!
 - Similarly, the next scheduler class of rt_sched_class should be wrr_sched_class
- Define necessary functions used for defining wrr_sched_class
 - enqueue_task, dequeue_task, pick_next_task, ...
 - You don't need to implement all those functions
- Define other necessary functions for load balancing or debugging

Modify kernel/sched/core.c to support WRR

- Problem: It assumes that there are only `rt_sched_class` and `fair_sched_class`
- We need to make sure that they are aware of `wrr_sched_class` too!
 - Initialize WRR run queue
 - Make `SCHED_WRR` policy valid
 - Manage forked tasks
 - the child task should follow the same scheduler policy of parent
 - ...

Debugging WRR scheduler

- Remind: You should turn on CONFIG_SCHED_DEBUG option in artik10_defconfig
- You might want to modify kernel/sched/debug.c to check whether your WRR scheduler works properly or not
- Scheduling information is written to /proc/sched_debug
- In print_...()
 - You can print additional informations here like scheduling policy, wrr_weight, ...
- In print_...()
 - Print additional statistics for WRR

Implement system calls

- You all know how to implement system calls!
- Authentication is important in sched_setweight()
 - Only the administrator may increase and decrease a process' weight
 - The user who owns the process may decrease its process' weight
 - Other users cannot adjust the process' weight
 - You can check the process' uid and euid to justify the ownership
- Nothing hard here :)

Experiment on WRR scheduler

- Main question: How the weight of WRR affects performance
- Measure the time for Trial program to finish for varying
 - Weights
 - Number of processes
 - ...
- Important: You should make sure that all 8 cores are active when you start your experiment!
 - Initially, it is highly likely to have only one core active
 - You can make some number (about 10) of processes run for some time to make all CPUs active

More things...

- CFS is highly optimized, while your scheduler is not → Slow!
 - When the shell is not responding using WRR, just wait for a while
 - It's worse when only one CPU is on
 - It takes some time for other CPUs become active...
 - Do not make many processes in once (ex: forking 100 processes)
- Write multicore-related code only when CONFIG_SMP option is on!
 - You can make a use of `#ifdef`
- It's safe to have `rcu_read_lock()` when you are iterating on CPU cores
- This project is harder than project 2, so start early!

About submission (IMPORTANT!)

- Make sure your branch name: *proj3*
- Don't be late!
 - TA will not grade the commits after the **deadline**.
- Slides and Demo
 - Send it to the TA's email (os-tas@spl.snu.ac.kr) before the **deadline**.
 - os-tas@spl.snu.ac.kr
 - Title: **[OS-ProjX] TeamX slides&demo submission**
 - File name: **TeamX-slides.ppt(.pdf), TeamX-demo.mp4(.avi....)**
- Check for format : slides title / demo name / branch name and directory name
- Please aggregate your demo videos (=submit only one video!)

Announcement

- Design Review
 - Team1~7 ⇒ Sungwoo Cho: pigbug419 **at** gmail **dot** com
 - Available schedule
 - Fri, 13:00~17:00
 - Tue, 10:00~15:25, 17:00~
 - Team8~14 ⇒ Kyungtae Kim: heaven **at** snu **dot** ac **dot** kr
- Check your source code before submission
 - There were some codes which were not compiled....

Q&A

Implement load balancing (1)

- Q1. How to check the remaining time slice or figure out when to trigger load balance?
- scheduler_tick()
 - in kernel/sched/core.c
 - Called every tick
- Tick frequency: HZ
 - A macro which represents the number of ticks in a second
 - For arm architecture, HZ = 100
 - **You shouldn't make an assumption that it is 100**

Implement load balancing (1)

- Q1. How to check the remaining time slice or figure out when to trigger load balance?
- scheduler_tick()
- Tick frequency: HZ
- jiffies
 - A global variable contains the number of ticks after system booting
 - unsigned long value - overflow could happen!
 - There are macros for comparing time
 - time_after(), time_before(), time_after_eq(), time_before_eq()
 - More things: <http://www.makelinux.net/books/lkd2/ch10lev1sec3>

Implement Load Balancing (2)

- Q2. How to determine the task is migratable or not?
- Tasks that are currently running are not eligible to be moved
- Some tasks may have some restrictions on which CPU they can be run on
 - How can we know it? and Why do they have restrictions?
 - Refer to existing load balancing codes to find the answer

Implement load balancing (3)

- Q3. How to prevent race condition while load balancing?
- scheduler_tick is called for every available CPU!
 - You need to make sure that only one thread is working on load balancing at any time!
- One seemingly simple & plausible solution
 - Make only a certain CPU can do load balancing
 - But, because CONFIG_HOTPLUG_CPU is on by default, the designated CPU could be turned off anytime...
 - What happens if the designated CPU is turned off? How can we prevent it?
- Think carefully about synchronization issues and hotplug CPU!