

Project1 & Kernel Programming

March 21, 2017

SPL, SNU

Project 1

General explanation about proj1

- Writing a system call
 - `int ptree(struct prinfo *buf, int *nr);`
 - It should be assigned number **380**
- Test your system call
 - Print the entire process tree in pre-order

Example program output

- Swapper(pid 0)
 - the process is used to represent the state of 'not working'
- Systemd(pid 1)
 - Manage all the process, init system.
- kthreadd(pid 2)
 - kernel thread daemon.
 - kthread_create()


```
systemd,1,1,0,156,2,0
systemd-journal,156,1,1,0,185,0
systemd-udevd,185,1,1,0,484,0
syslogd,484,1,1,0,495,0
...
deviced,802,1,1,1612,857,0
systemctl,1612,64,802,0,1613,0
systemctl,1613,64,802,0,1614,0
systemctl,1614,64,802,0,31175,0
...
kthreadd,2,1,0,3,0,0
ksoftirqd/0,3,1,2,0,5,0
kworker/0:0H,5,1,2,0,6,0
kworker/u8:0,6,1,2,0,7,0
```

Return Value

- `int ptree(struct prininfo *buf, int *nr);`
- Success
 - Your system call should return the total number of entries (this may be bigger than the actual number of entries copied).
 - `nr` can be changed
- Error
 - error handling: `-EINVAL` or `-EFAULT`
 - defined in `include/uapi/asm-generic/errno-base.h`

nr can be changed

- # of total processes: 100



Error Handling

- **-EINVAL**
 - if buf or nr are null, or the number of entries is less than 1
- **-EFAULT**
 - if buf or nr are **outside the accessible address space**

Error Handling

- How to print error message?
 - `int result = your_system_call(380,);`
 - `printf ("%d", result);`
 - ?????
- You cannot get -EINVAL or EFAULT on the return value
 - Use **errno** to print it out.
 - Use **perror()** function

Check before submission!

- **Unsafe access** to user-space memory
- Return value
 - Incorrect return value
 - Not modifying *nr when needed
- etc...

Remind

Concise README file

- Describe how to build your kernel
- Describe the high-level design and implementation
- Investigation of the process tree
- Any lessons learned

Project 1 – what to submit

- Concise 4-minute presentation slides (including your video demo) to ~~os staff~~ os-tas
 - **Limit: 10 slides including the title slide**
 - **I won't look at slides after the 10th.**
- Your presentation includes
 - A. High-level design and implementation
 - B. A video clip** that shows that your system works
 - C. The investigation of the process trees
 - D. Lessons learned

About submission (IMPORTANT!)

- Make sure your branch name: *proj1*
- Don't be late!
 - TA will not grade the commits after the **deadline**.
- Slides and Demo
 - Send it to the TA's email (os-tas@spl.snu.ac.kr) before the **deadline**.
 - os-tas@spl.snu.ac.kr
 - Title: [OS-ProjX] TeamX slides&demo submission
 - File name: TeamX-slides.ppt(.pdf), TeamX-demo.mp4(.avi....)

Kernel Programming

Directories you may need to take a look...

- **./arch**
 - Architecture-dependent (i.e. x86, arm, mips, ...) parts of linux
 - Tizen is **arm-based** system :)
- **./kernel**
 - Common kernel codes
- **./net**
 - Common network-related codes
- **./drivers**
 - Common driver codes for linux
- **./fs**
 - Common file system codes for linux
- **./include**
 - Common header files

Some things you should keep in mind...

- No memory protection
 - Corrupting kernel memory space will make the whole machine crash!
- No floating point or MMX operation
 - Real number calculation might be challenging and painful!
 - You might have to do it on some projects, though :)
- A rigid stack limit
 - Be cautious about allocating local arrays or having recursive calls
 - Use `kmalloc()` instead for huge arrays
- Your kernel code will be run in multi-core environment
 - Use proper synchronization mechanism to avoid race conditions
 - Be aware of deadlocks

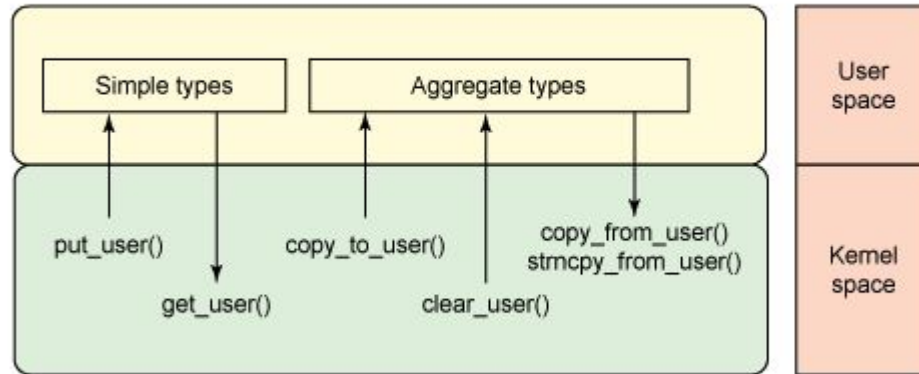
Accessing user-space memory

- In kernel mode, you should avoid accessing user memory space directly
 - Can result in kernel panic
- `include/asm/uaccess.h` provides user-memory-accessing macros
 - `get_user()` / `put_user()`: copies simple variables
 - `copy_from_user()` / `copy_to_user()`: copies a block of data
 - More things: <http://www.ibm.com/developerworks/library/l-kernel-memory-access/>
- Mark system call parameters containing user-space memory pointer with “__user”
 - e.g. In `./include/linux/syscalls.h`: `asmlinkage long sys_time(time_t __user *tloc);`

Function	Description
<code>access_ok</code>	Checks the validity of the user space memory pointer
<code>get_user</code>	Gets a simple variable from user space
<code>put_user</code>	Puts a simple variable to user space
<code>clear_user</code>	Clears, or zeros, a block in user space
<code>copy_to_user</code>	Copies a block of data from the kernel to user space
<code>copy_from_user</code>	Copies a block of data from user space to the kernel
<code>strnlen_user</code>	Gets the size of a string buffer in user space
<code>strncpy_from_user</code>	Copies a string from user space into the kernel

User Space Memory Access API

Figure 4. Data movement using the User Space Memory Access API



kmalloc() / kfree()

- Used for allocating / releasing kernel memory instead of malloc() / free()
 - Defined in linux/slab.h
- kmalloc() usage is similar to malloc(), but has an additional flag parameter
 - void *kmalloc(size_t size, int flags)
 - Frequently used flags
 - GFP_KERNEL: Allocate kernel-space memory
 - GFP_USER: Allocate user-space memory
 - GFP_ATOMIC: Similar to GFP_KERNEL, but allocation process cannot sleep. Used inside interrupts or other non-sleep routines.
 - More things: <http://www.makelinux.net/books/lkd2/ch11lev1sec4>
- kfree() usage is similar to free()

Task_struct

- About 300 lines!
- children, and sibling list → Doubly linked list

```
struct task_struct __rcu *real_parent; /* real parent process */
struct task_struct __rcu *parent; /* recipient of SIGCHLD, wait4() reports */
/*
 * children/sibling forms the list of my natural children
 */
struct list_head children; /* list of my children */
struct list_head sibling; /* linkage in my parent's children list */
```


Doubly linked list in linux kernel

- Linux kernel has a doubly linked list implementation for kernel programming
 - Extensively used across all linux kernel codes
- Defined in include/linux/list.h
 - Can only be used in kernel space!
- Unlike other commonly used linked lists, kernel list nodes are stored inside data

```
struct student {  
    char* name;  
    char* student_id;  
    struct list_head list;  
};
```

Doubly linked list in linux kernel (Contd.)

```
struct list_head {  
    struct list_head *next, *prev;  
}
```



Doubly linked list in linux kernel (Contd.)

- Defining list head pointer
 - `LIST_HEAD(student_list)`
- Initializing list node
 - `INIT_LIST_HEAD(&first_student->list);`
- More things on linux kernel list (highly recommended)
 - <http://www.makelinux.net/ldd3/chp-11-sect-5>
 - <https://bbingju.wordpress.com/2013/08/25/linux-kernel-list-h/> (Korean document)

Doubly linked list in linux kernel (Contd.)

- Commonly used macros/functions
 - `list_add()` / `list_add_tail()`: adding a node to a list
 - `list_del()` / `list_del_init()`: deleting a node from a list
 - `list_for_each_entry()`: iterating a list
 - `list_for_each_entry_safe()`: iterating a list **when nodes could be deleted**
 - and many more...(`for_each_task()`:??)

Task_struct - schedule info

- volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */

- ```
#define TASK_RUNNING 0
#define TASK_INTERRUPTIBLE 1
#define TASK_UNINTERRUPTIBLE 2
#define __TASK_STOPPED 4
#define __TASK_TRACED 8
/* in tsk->exit_state */
#define EXIT_ZOMBIE 16
#define EXIT_DEAD 32
/* in tsk->state again */
#define TASK_DEAD 64
#define TASK_WAKEKILL 128
#define TASK_WAKING 256
#define TASK_PARKED 512
#define TASK_STATE_MAX 1024
```

# Task\_struct - schedule info



- `int prio, static_prio, normal_prio;`  
`unsigned int rt_priority;`  
`const struct sched_class *sched_class;`  
`struct sched_entity se;`  
`struct sched_rt_entity rt;`
- Scheduling policy and priority

# Task\_struct - cpu mask

- `int nr_cpus_allowed;`  
`cpumask_t cpus_allowed;`
- The task can be running in the cpu in `cpus_allowed`

# Some useful references

- Linux cross reference
  - <http://lxr.free-electrons.com/source/?v=3.10>
- Unreliable Guide To Hacking The Linux Kernel by Rusty Russel
  - <http://kernelbook.sourceforge.net/kernel-hacking.pdf>



Q & A

# Back-up Slides